

Robust Hierarchical Clustering *

Maria-Florina Balcan, Yingyu Liang, Pramod Gupta
College of Computing, Georgia Institute of Technology
ninamf@cc.gatech.edu, {yliang39, pramodgupta}@gatech.edu

Abstract

One of the most widely used techniques for data clustering is agglomerative clustering. Such algorithms have been long used across many different fields ranging from computational biology to social sciences to computer vision in part because their output is easy to interpret. Unfortunately, it is well known, however, that many of the classic agglomerative clustering algorithms are not robust to noise (Narasimhan et al., 2006). In this paper we propose and analyze a new robust algorithm for bottom-up agglomerative clustering. We show that our algorithm can be used to cluster accurately in cases where the data satisfies a number of natural properties and where the traditional agglomerative algorithms fail. We also show how to adapt our algorithm to the inductive setting where our given data is only a small random sample of the entire data set. Experimental evaluations on synthetic and real world data sets show that our algorithm consistently achieves better performance than other hierarchical algorithms in the presence of noise.

Keywords: Unsupervised Learning, Clustering, Agglomerative Algorithms, Robustness

1 Introduction

Many data mining and machine learning applications ranging from computer vision to biology problems have recently faced an explosion of data. As a consequence it has become increasingly important to develop effective, accurate, robust to noise, fast, and general clustering algorithms, accessible to developers and researchers in a diverse range of areas.

One of the oldest and most commonly used methods for clustering data, widely used in many scientific applications, is hierarchical clustering (Gower, 1967; Bryant and Berry, 2001; Cheng et al., 2006; Dasgupta and Long, 2005; Duda et al., 2000; Gollapudi et al., 2006; Jain and Dubes, 1981; Jain et al., 1999; Johnson, 1967; Narasimhan et al., 2006). In hierarchical clustering the goal is not to find a single partitioning of the data, but a hierarchy (generally represented by a tree) of partitions which may reveal interesting structure in the data at multiple levels of granularity. The most widely used hierarchical methods are the agglomerative clustering techniques; most of these techniques start with a separate cluster for each point and then progressively merge the two closest clusters until only a single cluster remains. In all cases, we assume that we have a measure of similarity between pairs of objects, but the different schemes are distinguished by how they convert this into a measure of similarity between two clusters. For example, in single linkage the similarity between two clusters is the maximum similarity between points in these two different clusters. In complete linkage, the similarity between two clusters is the minimum similarity between points in these

*A preliminary version of this article appeared under the title *Robust Hierarchical Clustering* in the Proceedings of the Twenty-Third Conference on Learning Theory, 2010.

two different clusters. Average linkage defines the similarity between two clusters as the average similarity between points in these two different clusters (Dasgupta and Long, 2005; Jain et al., 1999).

Such algorithms have been used in a wide range of application domains ranging from biology applications to social sciences to computer vision applications mainly because they are quite fast and the output is quite easy to interpret. It is well known, however, that one of the main limitations of the agglomerative clustering algorithms is that they are not robust to noise (Narasimhan et al., 2006). In this paper we propose and analyze a robust algorithm for bottom-up agglomerative clustering. We show that our algorithm satisfies formal robustness guarantees and it will be successful in many cases where the traditional agglomerative algorithms fail.

In order to formally analyze correctness of our algorithm we use the framework introduced by Balcan et al. (2008). In this framework, we assume there is some target clustering (much like a k -class target function in the multi-class learning setting) and we say that an algorithm correctly clusters data satisfying property P if on any data set having property P , the algorithm produces a tree such that the target is some pruning of the tree. For example if all points are more similar to points in their own target cluster than to points in any other cluster (this is called the strict separation property), then any of the standard agglomerative algorithms will succeed. See Figure 1. However, with just tiny bit of noise, for example if each point has even just one point from a different cluster that it is similar to, then the standard algorithms will all fail (we elaborate on this in Section 2.2). See Figure 2. This brings up the question: is it possible to design an agglomerative algorithm that is robust to these types of situations and more generally can tolerate a substantial degree of noise? The contribution of our paper is to provide a strong positive answer to this question; we develop a robust, linkage based algorithm that will succeed in many interesting cases where standard agglomerative algorithms will fail. At a high level, our new algorithm is robust to noise in two different and important ways. First, it uses more global information for determining the similarities between clusters; second, it uses a robust linkage procedure involving a median test for linking the clusters, outvoting noisy similarities.

1.1 Our Results

In particular, in Section 3 we show that if the data satisfies a natural good neighborhood property, then our algorithm can be used to cluster well in the tree model, that is, to output a hierarchy such that the target clustering is (close to) a pruning of that hierarchy. The good neighborhood property relaxes the strict separation property, and only requires that after a small number of bad points (which could be extremely malicious) have been removed, for the remaining good points in the data set, in the neighborhood of their target cluster’s size, most of their nearest neighbors are from their target cluster. We show that our algorithm produces a hierarchy with a pruning that assigns all good points correctly. In Section 4, we further generalize this to allow for a good fraction of “boundary” points that do not fully satisfy the good neighborhood property. Unlike the good points, these points may have many nearest neighbors outside their target cluster in the neighborhood of their target cluster’s size; but also unlike the bad points, they have additional structure: they fall into a sufficiently large subset of their target cluster, such that all points in this subset have most of their nearest neighbors from this subset. As long as the fraction of boundary points in such subsets is not too large, our algorithm can produce a hierarchy with a pruning that assigns all good and boundary points correctly.

We further show how to adapt our algorithm to the inductive setting with formal correctness guarantees. In this setting, the clustering algorithm only uses a small random sample over the data set and generates a hierarchy over this sample, which also implicitly represents a hierarchy over the entire data set. This is especially useful when the amount of data is enormous such as in astrophysics and biology. We prove that

our algorithm requires only a small random sample whose size is independent of that of the entire data set and depends only on the noise and the confidence parameters.

We then perform experimental evaluations of our algorithm on synthetic data and real-world data sets. In the controlled experiments on synthetic data, our algorithm achieves results consistent with our theoretical analysis, outperforming several other hierarchical algorithms. Furthermore, in the experiments on real-world data, it performs consistently better than the other hierarchical algorithms. These experimental results suggest that the properties and the algorithm we propose can handle noise in real-world data as well.

1.2 Related Work

In agglomerative hierarchical clustering (Dasgupta and Long, 2005; Duda et al., 2000; Jain and Dubes, 1981; Jain et al., 1999), the goal is not to find a single partitioning of the data, but a hierarchy (generally represented by a tree) of partitionings which may reveal interesting structure in the data at multiple levels of granularity. Traditionally, only clusterings at a certain level are considered, but as we argue in Section 2 it is more desirable to consider all the prunings of the tree, since this way we can then handle much more general situations.

As mentioned above, it is well known that standard agglomerative hierarchical clustering techniques are not tolerant to noise (Nagy, 1968; Narasimhan et al., 2006). Several heuristic algorithms have been proposed to make the hierarchical clustering techniques more robust to noise, such as Ward’s method (Ward, 1963), Wishart’s method (Wishart, 1969), and CURE (Guha et al., 1998). However, these algorithms have no theoretical guarantees for their robustness. Also, our empirical study demonstrates that our algorithm has better tolerance to noise than these algorithms.

On the theoretical side, Balcan et al. (2008) analyzed the ν -strict separation property, a generalization of the simple strict separation property discussed above, requiring that after a small number of outliers have been removed all points are strictly more similar to points in their own cluster than to points in other clusters. They provided an algorithm for producing a hierarchy such that the target clustering is close to some pruning of the tree, but via a much more computationally expensive (non-agglomerative) algorithm. Our algorithm is simpler and substantially faster. As discussed in Section 2.1, the good neighborhood property is much broader than the ν -strict separation property, so our algorithm is much more generally applicable compared to their algorithm specifically designed for ν -strict separation.

In a different statistical model, Chaudhuri and Dasgupta (2010) proposed a generalization of Wishart’s method. They proved that given a sample from a density function, the method constructs a tree that is consistent with the cluster tree of the density. Although not directly targeting at robustness, the analysis shows the method successfully identifies salient clusters separated by low density regions, which suggests the method can be robust to the noise represented by the low density regions.

1.3 Structure of the Paper

The rest of the paper is organized as follows. In Section 2, we formalize our model and introduce the definition of the good neighborhood property. We describe our algorithm and prove it succeeds under the good neighborhood property in Section 3. We then prove that it also succeeds under a generalization of the good neighborhood property in Section 4. In Section 5, we show how to adapt our algorithm to the inductive setting with formal correctness guarantees. We provide the experimental results in Section 6, and conclude the paper in Section 7.

2 Definitions. A Formal Setup

We consider a clustering problem (S, ℓ) specified as follows. Assume we have a data set S of n objects. Each $x \in S$ has some (unknown) “ground-truth” label $\ell(x)$ in $Y = \{1, \dots, k\}$, where we will think of k as much smaller than n . Let $C_i = \{x \in S : \ell(x) = i\}$ denote the set of points of label i (which could be empty), and denote the target clustering as $\mathcal{C} = \{C_1, \dots, C_k\}$. Let $C(x)$ be a shorthand of $C_{\ell(x)}$, and n_C denote the size of a cluster C .

Given another proposed clustering h , $h : S \rightarrow Y$, we define the error of h with respect to the target clustering to be

$$\mathbf{err}(h) = \min_{\sigma \in \mathcal{S}_k} \left[\Pr_{x \in S} [\sigma(h(x)) \neq \ell(x)] \right] \quad (1)$$

where \mathcal{S}_k is the set of all permutations on $\{1, \dots, k\}$. Equivalently, the error of a clustering $\mathcal{C}' = \{C'_1, \dots, C'_k\}$ is $\min_{\sigma \in \mathcal{S}_k} \frac{1}{n} \sum_i |C_i - C'_{\sigma(i)}|$. This is popularly known as *Classification Error* (Meilă and Heckerman, 2001; Balcan et al., 2013; Voevodski et al., 2012).

We will be considering clustering algorithms whose only access to their data is via a pairwise similarity function $\mathcal{K}(x, x')$ that given two examples outputs a number in the range $[-1, 1]$. We will say that \mathcal{K} is a symmetric similarity function if $\mathcal{K}(x, x') = \mathcal{K}(x', x)$ for all x, x' . In this paper we assume that the similarity function \mathcal{K} is symmetric.

Our goal is to produce a hierarchical clustering that contains a pruning that is close to the target clustering. Formally, the goal of the algorithm is to produce a hierarchical clustering: that is, a tree on subsets such that the root is the set S , and the children of any node S' in the tree form a partition of S' . The requirement is that there must exist a *pruning* h of the tree (not necessarily using nodes all at the same level) that has error at most ϵ . Balcan et al. (2008) have shown that this type of output is necessary in order to be able to analyze non-trivial properties of the similarity function. For example, even if the similarity function satisfies the requirement that all points are more similar to all points in their own cluster than to any point in any other cluster (this is called the strict separation property) and even if we are told the number of clusters, there can still be multiple different clusterings that satisfy the property. In particular, one can show examples of similarity functions and two significantly different clusterings of the data satisfying the strict separation property. See Figure 1 for an example. However, under the strict separation property, there is a single hierarchical decomposition such that any consistent clustering is a pruning of this tree. This motivates clustering in the tree model, which is the model we consider in this work as well.

Given a similarity function satisfying the strict separation property (see Figure 1 for an example), we can efficiently construct a tree such that the ground-truth clustering is a pruning of this tree (Balcan et al., 2008). Moreover, almost any of the standard linkage based algorithms (for example, single linkage, average linkage, or complete linkage) would work well under this property. However, one can show that if the similarity function slightly deviates from the strict separation condition, then all the standard agglomerative algorithms will fail (we elaborate on this in section 2.2). In this context, the main question we address in this work is: Can we develop other more robust, linkage based algorithms that will succeed under more realistic and yet natural conditions on the similarity function?

Note: Note that strict separation does not guarantee that all the cutoffs for different points x are the same, so single linkage would not necessarily have the right clustering if it just stopped once it has k clusters; however the target clustering will provably be a pruning of the final single linkage tree; this is why we define success based on prunings.

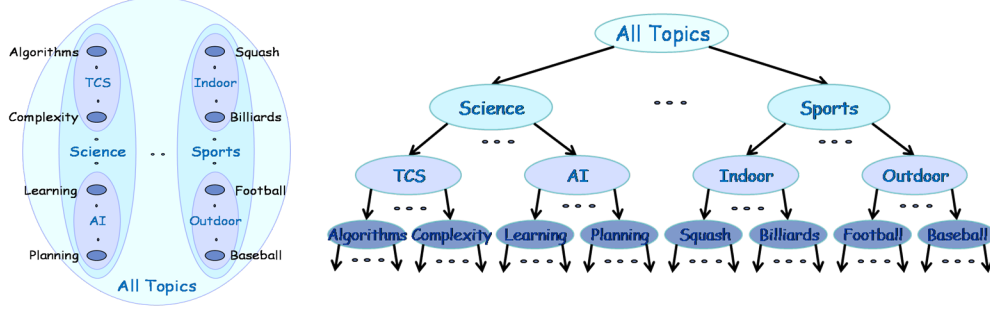


Figure 1: Consider a document clustering problem. Assume that data lies in multiple regions Algorithms, Complexity, Learning, Planning, Squash, Billiards, Football, Baseball. Suppose that the similarity $\mathcal{K}(x, y) = 0.999$ if x and y belong to the same inner region; $\mathcal{K}(x, y) = 3/4$ if $x \in \text{Algorithms}$ and $y \in \text{Complexity}$, or if $x \in \text{Learning}$ and $y \in \text{Planning}$, or if $x \in \text{Squash}$ and $y \in \text{Billiards}$, or if $x \in \text{Football}$ and $y \in \text{Baseball}$; $\mathcal{K}(x, y) = 1/2$ if x is in (Algorithms or Complexity) and y is in (Learning or Planning), or if x is in (Squash or Billiards) and y is in (Football or Baseball); define $\mathcal{K}(x, y) = 0$ otherwise. Both clusterings $\{\text{Algorithms} \cup \text{Complexity} \cup \text{Learning} \cup \text{Planning}, \text{Squash} \cup \text{Billiards}, \text{Football} \cup \text{Baseball}\}$ and $\{\text{Algorithms} \cup \text{Complexity}, \text{Learning} \cup \text{Planning}, \text{Squash} \cup \text{Billiards} \cup \text{Football} \cup \text{Baseball}\}$ satisfy the strict separation property.

2.1 Properties of the Similarity Function

We describe here some natural properties of the similarity functions that we analyze in this paper. We start with a noisy version of the simple strict separation property (mentioned above) which was introduced in (Balcan et al., 2008) and we then define an interesting and natural generalization of it.

Property 1. *The similarity function \mathcal{K} satisfies ν -strict separation for the clustering problem (S, ℓ) if for some $S' \subseteq S$ of size $(1 - \nu)n$, \mathcal{K} satisfies strict separation for (S', ℓ) . That is, for all $x, x', x'' \in S'$ with $x' \in C(x)$ and $x'' \notin C(x)$ we have $\mathcal{K}(x, x') > \mathcal{K}(x, x'')$.*

So, in other words we require that the strict separation is satisfied after a number of bad points have been removed. A somewhat different condition is to allow each point to have some bad immediate neighbors as long as most of its immediate neighbors are good. Formally:

Property 2. *The similarity function \mathcal{K} satisfies α -good neighborhood property for the clustering problem (S, ℓ) if for all points x we have that all but αn out of their $n_{C(x)}$ nearest neighbors belong to the cluster $C(x)$.*

Note that α -good neighborhood property is different from the ν -strict separation property. For the ν -strict separation property we can have up to νn bad points that can misbehave; in particular, these νn bad points can have similarity 1 to *all* the points in S ; however, once we remove these points the remaining points are more similar to points in their own cluster than to points in other cluster. On the other hand, for the α -good neighborhood property we require that for all points x all but αn out of their $n_{C(x)}$ nearest neighbors belong to the cluster $C(x)$. (So we cannot have a point that has similarity 1 to all the points in S .) Note however that different points might misbehave on different αn neighbors. We can also consider a property that generalizes both the ν -strict separation property and the α -good neighborhood. Specifically:

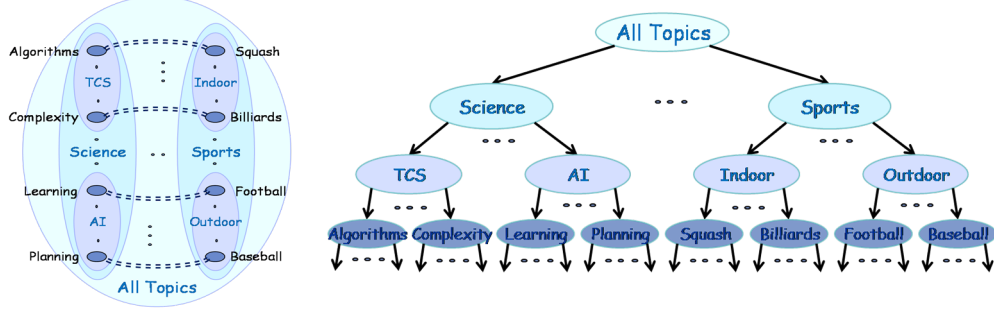


Figure 2: Same as Figure 1 except that let us match each point in Algorithms with a point in Squash, each point in Complexity with a point in Billiards, each point in Learning with a point in Football, and each point in Planning with a point in region Baseball. Define the similarity measure to be the same as in Figure 1 except that we let $\mathcal{K}(x, y) = 1$ if x and y are matched. Note that for $\alpha = 1/n$ the similarity function satisfies the α -good neighborhood with respect to any of the prunings of the tree above. However, single linkage, average linkage, and complete linkage would initially link the matched pairs and produce clusters with very high error with respect to any such clustering.

Property 3. *The similarity function \mathcal{K} satisfies (α, ν) -good neighborhood property for the clustering problem (S, ℓ) if for some $S' \subseteq S$ of size $(1 - \nu)n$, \mathcal{K} satisfies α -good neighborhood for (S', ℓ) . That is, for all points $x \in S'$ we have that all but αn out of their $n_{C(x) \cap S'}$ nearest neighbors in S' belong to the cluster $C(x)$.*

Clearly, (α, ν) -good neighborhood is a generalization of both ν -strict separation and α -good neighborhood. Formally,

Fact 1. *If the similarity function \mathcal{K} satisfies the α -good neighborhood property for the clustering problem (S, ℓ) , then \mathcal{K} also satisfies the $(\alpha, 0)$ -good neighborhood property for the clustering problem (S, ℓ) .*

Fact 2. *If the similarity function \mathcal{K} satisfies the ν -strict separation property for the clustering problem (S, ℓ) , then \mathcal{K} also satisfies the $(0, \nu)$ -good neighborhood property for the clustering problem (S, ℓ) .*

Our algorithm also succeeds under a more general property that we call weak good neighborhood property, which allows a good fraction of points to not fully satisfy the good neighborhood property. For clarity of exposition, we do not introduce the definition here, but discuss it in detail in Section 4.

Balcan et al. (2008) have shown that if \mathcal{K} satisfies the ν -strict separation property with respect to the target clustering, then as long as the smallest target cluster has size $5\nu n$, one can in polynomial time construct a hierarchy such that the ground-truth is ν -close to a pruning of the hierarchy. Unfortunately the algorithm presented in (Balcan et al., 2008) is computationally very expensive: it first generate a large list of $\Omega(n^2)$ candidate clusters and repeatedly runs pairwise tests in order to laminarize these clusters; its running time is a large unspecified polynomial. Our new robust linkage algorithm can be used to get a simpler and much faster algorithm for clustering accurately under the ν -strict separation property. Additionally, our algorithm succeeds under the weak good neighborhood property, which is much more general than ν -strict separation.

2.2 Standard Linkage Based Algorithms Are Not Robust

As we show below, even if the data satisfies the good neighborhood property, still most of the standard linkage based algorithms will fail. The contribution of our work is to develop robust, linkage based algorithms

that will succeed under these natural conditions. More specifically, we can show an example where simple linkage based algorithm would perform very badly, but where our algorithm would work well. In particular, if we slightly modify the example in Figure 1, by adding a little bit of noise, to form links of high similarity between points in different inner blobs, we can show that many of the classic linkage based algorithms will perform poorly¹. See Figure 2 for a precise description of the similarity.

In particular, the single linkage algorithm, the average linkage algorithm, and the complete linkage algorithm, would in the first $n/2$ stages merge the matched pairs of points. From that moment on, no matter how they perform, none of the natural and desired clusterings will even be $1/2$ close to any of the prunings of the hierarchy produced. Notice however, that \mathcal{K} satisfies α -good neighborhood with respect to any of the desired clusterings (for $\alpha = 1/n$), and that our algorithm will be successful on this instance. The ν -strict separation is not satisfied in this example either, for any constant ν .

3 Robust Median Neighborhood Linkage

In this section, we propose a new algorithm, Robust Median Neighborhood Linkage, and show that it succeeds for instances satisfying (α, ν) -good neighborhood property.

Informally, the algorithm maintains a threshold t and a list \mathcal{C}'_t of subsets of points of S ; these subsets are called blobs for convenience. We first initialize the threshold to a value t that is not too large and not too small ($t = 6(\alpha + \nu)n + 1$), and initialize \mathcal{C}'_{t-1} to contain $|S|$ blobs, one for each point in S . For each t , the algorithm builds \mathcal{C}'_t from \mathcal{C}'_{t-1} by merging two or more blobs as follows. It first builds a graph F_t , whose vertices are the data points in S and whose edges are constructed by connecting any two points that share at least $t - 2(\alpha + \nu)n$ points in common out of their t nearest neighbors. Then it builds a graph H_t whose vertices correspond to blobs in \mathcal{C}'_{t-1} and whose edges are specified in the following way. Two singleton blobs $C_u = \{x\}$ and $C_v = \{y\}$ are connected in H_t if the points x, y have more than $(\alpha + \nu)n$ common neighbors in F_t . For blobs C_u and C_v that are not both singleton, the algorithm performs a median test. In this test, for each pair of points $x \in C_u, y \in C_v$, it computes the number $S_t(x, y)$ of points $z \in C_u \cup C_v$ that are the common neighbors of x and y in F_t . It then connects C_u and C_v in H_t if $\text{median}_{x \in C_u, y \in C_v} S_t(x, y)$ is larger than $1/4$ fraction of $|C_u| + |C_v|$. Once H_t is built, we analyze its connected components in order to create \mathcal{C}'_t . For each component V of H_t , if V contains sufficiently many points from S in its blobs we merge all its blobs into one blob in \mathcal{C}'_t . After building \mathcal{C}'_t , the threshold is increased and the above steps are repeated until only one blob is left. The algorithm finally outputs the tree with single points as leaves and internal nodes corresponding to the merges performed. The full details of our algorithm are described in Algorithm 1. Our main result in this section is the following:

Theorem 1. *Let \mathcal{K} be a symmetric similarity function satisfying the (α, ν) -good neighborhood for the clustering problem (S, ℓ) . As long as the smallest target cluster has size greater than $6(\nu + \alpha)n$, Algorithm 1 outputs a hierarchy such that a pruning of the hierarchy is ν -close to the target clustering in time $O(n^{\omega+1})$, where $O(n^\omega)$ is the state of the art for matrix multiplication.*

In the rest of this section, we will first describe the intuition behind the algorithm in Section 3.1 and then prove Theorem 1 in Section 3.2.

¹Since, usually, the similarity function between pairs of objects is constructed based on heuristics, this could happen in practice; for example we could have a similarity measure that puts a lot of weight on features such as date or names, and so we could easily have a document about Learning being more similar to a document about Football than to other documents about Learning. While this example seems a little bit contrived, in Figure 7 in Section 4 we will give a naturally occurring example where the linkage algorithms still fail but our algorithm succeeds because it satisfies a generalization of the good neighborhood property that we will discuss in Section 4.

Algorithm 1 Robust Median Neighborhood Linkage

Input: Similarity function \mathcal{K} on a set of points S , $n = |S|$, parameters $\alpha > 0, \nu > 0$.

Step 1 Initialize $t = 6(\alpha + \nu)n + 1$.
Initialize \mathcal{C}'_{t-1} to be a list of blobs so that each point is in its own blob.
while $|\mathcal{C}'_{t-1}| > 1$ **do**

Step 2 \triangleright Build a graph F_t whose vertices are points in S and
 whose edges are specified as follows.
 Let $N_t(x)$ denote the t nearest neighbors of x .
 for any $x, y \in S$ that satisfy $|N_t(x) \cap N_t(y)| \geq t - 2(\alpha + \nu)n$ **do**
 Connect x, y in F_t .
 end for

Step 3 \triangleright Build a graph H_t whose vertices are blobs in \mathcal{C}'_{t-1} and
 whose edges are specified as follows.
 Let $N_F(x)$ denote the neighbors of x in F_t .
 for any $C_u, C_v \in \mathcal{C}'_{t-1}$ **do**
 if C_u, C_v are singleton blobs, i.e. $C_u = \{x\}, C_v = \{y\}$ **then**
 Connect C_u, C_v in H_t , if $|N_F(x) \cap N_F(y)| > (\alpha + \nu)n$.
 else
 Set $S_t(x, y) = |N_F(x) \cap N_F(y) \cap (C_u \cup C_v)|$, i.e. the number of
 points in $C_u \cup C_v$ that are common neighbors of x, y in F_t .
 Connect C_u, C_v in H_t , if $\text{median}_{x \in C_u, y \in C_v} S_t(x, y) > \frac{|C_u| + |C_v|}{4}$.
 end if
 end for

Step 4 \triangleright Merge blobs in \mathcal{C}'_{t-1} to get \mathcal{C}'_t
 Set $\mathcal{C}'_t = \mathcal{C}'_{t-1}$.
 for any component V in H_t that satisfies $|\bigcup_{C \in V} C| \geq 4(\alpha + \nu)n$ **do**
 Update \mathcal{C}'_t by merging all blobs in V into one blob.
 end for

Step 5 \triangleright Increase threshold
 $t = t + 1$.
end while

Output: Tree T with single points as leaves and internal nodes corresponding to the merges performed.

3.1 Intuition of the Algorithm under the Good Neighborhood Property

We begin with some convenient terminology and a simple fact about the good neighborhood property. In the definition of (α, ν) -good neighborhood property (see Property 3), we call the points in S' good points and the points in $B = S \setminus S'$ bad points. Let $G_i = C_i \cap S'$ be the good set of label i . Let $G = \bigcup_i G_i$ denote the whole set of good points; so $G = S'$. Clearly $|G| \geq n - \nu n$. Recall that n_{C_i} is the number of points in the cluster C_i . Note that the following is a useful consequence of the (α, ν) -good neighborhood property.

Fact 3. *Suppose the similarity function \mathcal{K} satisfies the (α, ν) -good neighborhood property for the clustering problem (S, ℓ) . As long as t is smaller than n_{C_i} , for any good point $x \in C_i$, all but at most $(\alpha + \nu)n$ out of its t nearest neighbors lie in its good set G_i .*




Figure 3: Graph F and H when all target clusters are of the same size n_C , which is known. In F , no two good points in different target clusters can be connected, and all good points in the same target cluster will be connected. In H , bad points connected to good points from different target clusters are disconnected.

Proof. Let $x \in G_i$. By definition, out of its $|G_i|$ nearest neighbors in G , there are at least $|G_i| - \alpha n$ points from G_i . These points must be among its $|G_i| + \nu n$ nearest neighbors in S , since there are at most νn bad points in $S \setminus G$. This means that at most $(\alpha + \nu)n$ out of its $|G_i| + \nu n$ nearest neighbors are outside G_i . Notice $|G_i| + \nu n \geq n_{C_i}$, we have that at most $(\alpha + \nu)n$ out of its n_{C_i} nearest neighbors are outside G_i , as desired. \square

Intuition We first assume for simplicity that all the target clusters have the same size n_C and that we know n_C . In this case it is quite easy to recover the target clusters as follows. We first construct a graph F whose vertices are points in S ; we connect two points in F if they share at least $n_C - 2(\nu + \alpha)n$ points in common among their n_C nearest neighbors. By Fact 3, if the target clusters are not too small (namely $n_C > 6(\nu + \alpha)n$), we are guaranteed that no two good points in different target clusters will be connected in F , and that all good points in the same target cluster will be connected in F . If there are no bad points ($\nu = 0$), then each connected component of F corresponds to a target cluster, and we could simply output them. Alternatively, if there are bad points ($\nu > 0$), we can still cluster well as follows. We construct a new graph H on points in S by connecting points that share more than $(\alpha + \nu)n$ neighbors in the graph F . The key point is that in F a bad point can be connected to good points from only one single target cluster. This then ensures that no good points from different target clusters are in the same connected component in H . So, if we output the largest k components of H , we will obtain a clustering with error at most νn . See Figure 3 for an illustration.

If we do not know n_C , we can still use a pretty simple procedure. Specifically, we start with a threshold $t \leq n_C$ that is not too small and not too large (say $6(\nu + \alpha)n < t \leq n_C$), and build a graph F_t on S by connecting two points if they share at least $t - 2(\nu + \alpha)n$ points in common out of their t nearest neighbors. We then build another graph H_t on S by connecting points if they share more than $(\alpha + \nu)n$ neighbors in the graph F_t . The key idea is that when $t \leq n_C$, good points from different target clusters share less than $t - 2(\nu + \alpha)n$ neighbors, and thus are not connected in F_t and H_t . If the k largest connected components of H_t all have sizes greater than $(\alpha + \nu)n$ and they cover at least a $(1 - \nu)$ fraction of the whole set of points S , then these components must correspond to the target clusters and we can output them. Otherwise, we increase the critical threshold and repeat. By the time we reach n_C , all good points in the same target clusters will get connected in F_t and H_t , so we can identify the k largest components as the target clusters.

Note that as mentioned above, when $t \leq n_C$, each connected component in H_t can contain good points from only one target cluster. An alternative procedure is to reuse this information in later thresholds, so that we do not need to build the graph H_t from scratch as described in the above paragraph. Specifically, we




Figure 4: Graph F_t and H_t when target clusters have the same size n_C but we do not know n_C . The figure shows the case when $t < n_C$. In F_t , no good points are connected with good points outside their target clusters; in H_t , blobs containing good points in different target clusters are disconnected.

maintain a list \mathcal{C}'_t of subsets of points; these subsets are called blobs for convenience. We start with a list where each blob contains a single point. At each threshold t , we build F_t on the points in S as before, but build H_t on the blobs in \mathcal{C}'_{t-1} (instead of on the points in S) by connecting two blobs if they have points that share enough neighbors in F_t . Any two connected blobs must contain only good points from the same target cluster, so we can update the list of blobs by merging the connected blobs into one blob. See Figure 4 for an illustration of F_t and H_t .

Additionally, when building H_t , we can use a median test for non-singleton blobs to outvote the noise. In particular, for two blobs C_u and C_v that are not both singleton, we compute for all $x \in C_u$ and $y \in C_v$ the quantity $S_t(x, y)$, which is the number of points in $C_u \cup C_v$ that are the common neighbors of x and y in F_t . We then connect the two blobs in H_t if $\text{median}_{x \in C_u, y \in C_v} S_t(x, y)$ is sufficiently large. See Step 3 in Algorithm 1 for the details. This allows some points in these blobs to have weaker properties than the good neighborhood, leading to better robustness of our algorithm. See Section 4 for examples of such points and a theoretical analysis of the robustness.

In the general case where the sizes of the target clusters are different, similar ideas can be applied. The key point is that when $t \leq n_{C_i}$, good points from C_i share less than $t - 2(\nu + \alpha)n$ neighbors with good points outside, and thus are not connected to them in F_t . Then in H_t , we can make sure that no blobs containing good points in C_i will be connected with blobs containing good points outside C_i . When $t = n_{C_i}$, good points in C_i form a clique in F_t , then all the blobs containing good points in C_i are connected in H_t , and thus are merged. See Figure 5 for an illustration. Full details are presented in Algorithm 1 and the proof of Theorem 1 in the following subsection.

3.2 Correctness under the Good Neighborhood Property

In this subsection, we prove Theorem 1 for our algorithm. The correctness follows from Lemma 1 and the running time follows from Lemma 2. Before proving these lemmas, we begin with a useful fact which follows immediately from the design of the algorithm.

Fact 4. *In Algorithm 1, for any t , if a blob in \mathcal{C}_t contains at least one good point, then at least $3/4$ fraction of the points in that blob are good points.*

Proof. This is clearly true when the blob is singleton. When it is non-singleton, it must be formed in Step 4 in Algorithm 1, so it contains at least $4(\alpha + \nu)n$ points. Then the claim follows since there are at most νn bad points. \square




Figure 5: Graph F_t and H_t when target clusters are of different sizes. The figure shows the case when $t = n_{C_2}$. In F_t , no good points are connected with good points outside their target clusters; good points in C_2 form a clique since $t = n_{C_2}$. In H_t , blobs containing good points in different target clusters are disconnected; blobs containing good points in C_2 are all connected.

Lemma 1. *The following claims are true in Algorithm 1:*

- (1) *For any C_i such that $t \leq |C_i|$, any blob in \mathcal{C}'_t containing good points from C_i will not contain good points outside C_i .*
- (2) *For any C_i such that $t = |C_i|$, all good points in C_i belong to one blob in \mathcal{C}'_t .*

Proof. Before proving the claims, we first show that the graph F_t constructed in Step 2 has the following useful properties. Recall that F_t is constructed on points in S by connecting any two points that share at least $t - 2(\alpha + \nu)n$ points in common out of their t nearest neighbors. For any C_i such that $t \leq |C_i|$, we have:

- (a) If x is a good point in C_i and y is a good point outside C_i , then x and y are not connected in F_t .

To see this, first note that by Fact 3, x has at most $(\alpha + \nu)n$ neighbors outside C_i out of the t nearest neighbors. For $y \in G_j$, if $n_{C_j} \geq t$, then y has at most $(\alpha + \nu)n$ neighbors in C_i ; if $n_{C_j} < t$, y has at most $(\alpha + \nu)n + t - n_{C_j}$ neighbors in C_i . In both cases, y has at most $(\alpha + \nu)n + \max(0, t - n_{C_j}) < t - 5(\alpha + \nu)n$ neighbors in C_i , since $n_{C_j} > 6(\alpha + \nu)n$ and $t > 6(\alpha + \nu)n$. Then x and y have at most $t - 4(\alpha + \nu)n$ common neighbors, so they are not connected in F_t .

- (b) If x is a good point in C_i , y is a good point outside C_i , and z is a bad point, then z cannot be connected to both x and y in F_t .

To prove this, we will show that if z is connected to x , then z cannot be connected to y . First, by the same argument as above, out of the t nearest neighbors, y has less than $t - 5(\alpha + \nu)n$ neighbors in C_i . Second, by Fact 3, x has at most $(\alpha + \nu)n$ neighbors outside C_i . If z has less than $t - 3(\alpha + \nu)n$ neighbors in C_i , then z and x share less than $t - 3(\alpha + \nu)n + (\alpha + \nu)n = t - 2(\alpha + \nu)n$ neighbors and will not be connected. So z must have at least $t - 3(\alpha + \nu)n$ neighbors in C_i , and thus cannot have more than $3(\alpha + \nu)n$ neighbors outside C_i . The two statements show that y and z share less than $t - 5(\alpha + \nu)n$ neighbors in C_i , and at most $3(\alpha + \nu)n$ neighbors outside C_i . So they share less than $t - 2(\alpha + \nu)n + 3(\alpha + \nu)n = t - 2(\alpha + \nu)n$ neighbors and thus are not connected in F_t .

Now we prove Claim (1) in the lemma by induction on t . The claim is clearly true initially. Assume for induction that the claim is true for the threshold $t - 1 < |C_i|$, that is, for any C_i such that $t - 1 < |C_i|$, any blob in \mathcal{C}'_{t-1} containing good points from C_i will not contain good points outside C_i . We now prove that




Figure 6: Illustration of the median test on C_u and C_v . At least $3/4$ fraction of the points in C_u and C_v are good points, and thus more than half of the pairs (x, y) with $x \in C_u$ and $y \in C_v$ are pairs of good points.

the graph H_t constructed in Step 3 has the following properties, which can be used to show that the claim is still true for the threshold t .

- If $C_u \in \mathcal{C}'_{t-1}$ contains good points from C_i and $C_v \in \mathcal{C}'_{t-1}$ contains good points outside C_i , then they cannot be connected in H_t .

If both C_u and C_v are singleton blobs, say $C_u = \{x\}$, $C_v = \{y\}$, then by Property (a) of F_t , the common neighbors of x and y can only be bad points, and thus C_u and C_v cannot be connected.

If one of the two blobs (say C_u) is a singleton blob and the other is not, then C_u contains only one good point, and by Fact 4, at least $3/4$ fraction of the points in C_v are good points. If both C_u and C_v are non-singleton blobs, then by Fact 4, at least $3/4$ fraction of the points in C_u and C_v are good points. Therefore, in both cases, the number of pairs (x, y) with good points $x \in C_u$ and $y \in C_v$ is at least $\frac{3}{4}|C_u| \times \frac{3}{4}|C_v| > \frac{|C_u||C_v|}{2}$. That is, more than half of the pairs (x, y) with $x \in C_u$ and $y \in C_v$ are pairs of good points; see Figure 6 for an illustration. This means there exist good points $x^* \in C_u, y^* \in C_v$ such that $S_t(x^*, y^*)$ is no less than $\text{median}_{x \in C_u, y \in C_v} S_t(x, y)$. By the induction assumption, x^* is a good point in C_i and y^* is a good point outside C_i . Then by Property (a)(b) of F_t , x^* and y^* have no common neighbors in F_t , and thus $\text{median}_{x \in C_u, y \in C_v} S_t(x, y) = 0$. Therefore, C_u and C_v are not connected in H_t .

- If $C_u \in \mathcal{C}'_{t-1}$ contains good points from C_i , $C_v \in \mathcal{C}'_{t-1}$ contains good points outside C_i , and $C_w \in \mathcal{C}'_{t-1}$ contains only bad points, then C_w cannot be connected to both C_u and C_v in H_t .

To prove this, assume for contradiction that C_w is connected to both C_u and C_v . First, note the following fact about C_w . Since any non-singleton blob must be formed in Step 4 in the algorithm and contain at least $4(\alpha + \nu)n$ points and thus cannot contain only bad points, C_w must be a singleton blob, containing only a bad point z .

Next, we show that if $C_w = \{z\}$ is connected to C_u , then z must be connected to some good point in C_i in F_t . If C_u is a singleton blob, say $C_u = \{x\}$, then by Step 4 in the algorithm, z and x share more than $(\alpha + \nu)n$ common neighbors in F_t . By Property (a)(b) of F_t , the common neighbors of x and z in F_t can only be good points in C_i or bad points. Since there are at most νn bad points, z must be connected to some good point in C_i in F_t . If C_u is not a singleton blob, then by Step 4 in the algorithm, $\text{median}_{x \in C_u} S_t(x, z) > (|C_u| + |C_w|)/4$. By Fact 4, at least $3/4$ fraction of the points in C_u are good points. So there exists a good point $x^* \in C_u$ such that $S_t(x^*, z) \geq \text{median}_{x \in C_u} S_t(x, z)$, which leads to $S_t(x^*, z) > (|C_u| + |C_w|)/4 > \nu n$. By the induction assumption, x^* is a good point in C_i . Then by Property (a) of F_t , x^* is only connected to good points from C_i and bad points. Since $S_t(x^*, z) > \nu n$, z and x^* must share some common neighbors from C_i . Therefore, z is connected to

some good point in C_i in F_t .

Similarly, if $C_w = \{z\}$ is connected to C_v , z must be connected to some good point outside C_i in F_t . But then z is connected to both a good point in C_i and a good point outside F_t , which contradicts Property (b) of F_t .

By the properties of H_t , no connected component contains both good points in C_i and good points outside C_i . So Claim (1) is still true for the threshold t . By induction, it is true for all thresholds.

Finally, we prove Claim (2). First, at the threshold $t = |C_i|$, all good points in C_i are connected in F_t . This is because any good point in C_i has at most $(\alpha + \nu)n$ neighbors outside C_i , so when $t = |C_i|$, any two good points in C_i share at least $t - 2(\alpha + \nu)n$ common neighbors and thus are connected in F_t .

Second, all blobs in \mathcal{C}'_{t-1} containing good points in C_i are connected in H_t . There are two cases.

- If no good points in C_i have been merged, then all singleton blobs containing good points in C_i will be connected in H_t .

This is because all good points in C_i are connected in F_t , and thus they share at least $|G_i| \geq 6(\alpha + \nu)n - \nu n$ points as common neighbors in F_t .

- If some good points in C_i have already been merged into non-singleton blobs in \mathcal{C}'_{t-1} , we can show that in H_t these non-singleton blobs will be connected to each other and connected to singleton blobs containing good points from C_i .

Consider two non-singleton blobs C_u and C_v that contain good points from C_i . By Fact 4, at least $3/4$ fraction of the points in C_u and C_v are good points. So there exist good points $x^* \in C_u$ and $y^* \in C_v$ such that $S_t(x^*, y^*) \leq \text{median}_{x \in C_u, y \in C_v} S_t(x, y)$. By Claim (1), x^* and y^* must be good points in C_i . Then they are connected to all good points in C_i in F_t , and thus $S_t(x^*, y^*)$ is no less than the number of good points in C_u and C_v , which is at least $3(|C_u| + |C_v|)/4$. Now we have $\text{median}_{x \in C_u, y \in C_v} S_t(x, y) \geq S_t(x^*, y^*) \geq 3(|C_u| + |C_v|)/4 > (|C_u| + |C_v|)/4$, and thus C_u, C_v are connected in H_t .

Consider a non-singleton blob C_u and a singleton blob C_v that contain good points from C_i . The above argument also holds, so C_u, C_v are connected in H_t .

Therefore, in both cases, all blobs in \mathcal{C}'_{t-1} containing good points in C_i are connected in H_t . Then in Step 4, all good points in C_i are merged into a blob in \mathcal{C}'_t . \square

Lemma 2. *Algorithm 1 has a running time of $O(n^{\omega+1})$.*

Proof. The initializations in Step 1 take $O(n)$ time. To compute F_t in Step 2, for any $x \in S$, let $I_t(x, y) = 1$ if y is within the t nearest neighbors of x , and let $I_t(x, y) = 0$ otherwise. Initializing I_t takes $O(n^2)$ time. Next we compute $N_t(x, y)$, the number of common neighbors between x and y . Notice that $N_t(x, y) = \sum_{z \in S} I_t(x, z)I_t(y, z)$, so $N_t = I_t I_t^T$. Then we can compute the adjacency matrix F_t (overloading notation for the graph F_t) from N_t . These steps take $O(n^\omega)$ time.

To compute the graph H_t in Step 3, first define $NS_t = F_t(F_t)^T$. Then for two points x and y , $NS_t(x, y)$ is the number of their common neighbors in F_t . Further define matrix FC_t as follows: if x and y are connected in F_t and are in the same blob in \mathcal{C}'_{t-1} , then let $FC_t(x, y) = 1$; otherwise, let $FC_t(x, y) = 0$. As a reminder, for two points x that belongs to $C_u \in \mathcal{C}'_{t-1}$ and y that belongs to $C_v \in \mathcal{C}'_{t-1}$, $S_t(x, y)$ is the number of points in $C_u \cup C_v$ they share as neighbors in common in F_t . FC_t is useful for computing S_t :

since for $x \in C_u$ and $y \in C_v$,

$$\begin{aligned} S_t(x, y) &= \sum_{z \in C_v} F_t(x, z)F_t(y, z) + \sum_{z \in C_u} F_t(x, z)F_t(y, z) \\ &= \sum_{z \in S} F_t(x, z)FC_t(y, z) + \sum_{z \in S} FC_t(x, z)F_t(y, z), \end{aligned}$$

we have $S_t = F_t(FC_t)^T + FC_t(F_t)^T$. Based on NS_t and S_t , we can then build the graph H_t . All these steps take $O(n^\omega)$ time.

When we perform merges in Step 4 or increase the threshold in Step 5, we need to recompute the above data structures, which takes $O(n^\omega)$ time. Since there are $O(n)$ merges and $O(n)$ thresholds, Algorithm 1 takes time $O(n^{\omega+1})$ in total. \square

4 A More General Property: Weak Good Neighborhood

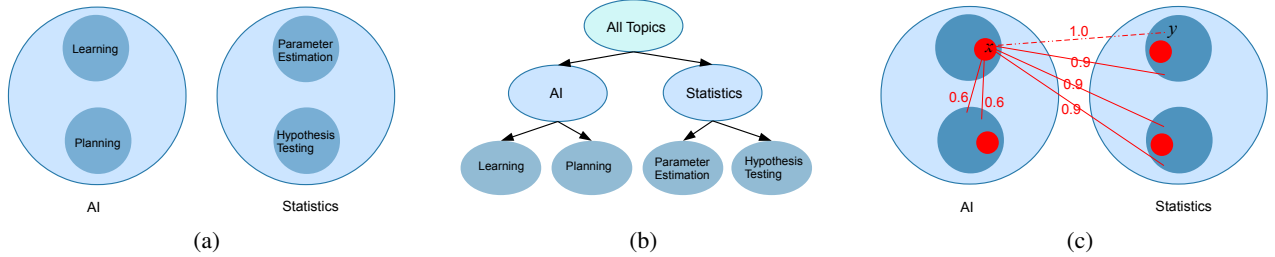


Figure 7: Consider a document clustering problem. Assume that there are $n/4$ documents in each of the four areas: Learning, Planning, ParameterEstimation and HypothesisTesting. The first two belong to the field AI, and the last two belong to the field Statistics. The similarities are specified as follows. (1) $\mathcal{K}(x, y) = 0.99$ if x, y belong to the same area; (2) $\mathcal{K}(x, y) = 0.8$ if x, y belong to different areas in the same field; (3) $\mathcal{K}(x, y) = 0.5$ if x, y belong to different fields. As shown in (b), there are four prunings: $\{\text{Learning, Planning, ParameterEstimation, HypothesisTesting}\}$, $\{\text{AI, ParameterEstimation, HypothesisTesting}\}$, $\{\text{Learning, Planning, Statistics}\}$, and $\{\text{AI, Statistics}\}$. All these four prunings satisfy the strict separation property, and consequently satisfy the α -good neighborhood property for $\alpha = 0$. However, this is no longer true if we take into account noise that naturally arises in practice. As shown in (c), in each area, $1/8$ fraction of the documents lie close to the boundary between the two fields. More precisely, the similarities for these boundary documents are defined as follows. (1) These documents are very similar to some document in the other field: for each boundary document x , we randomly pick one document y in the other field and set $\mathcal{K}(x, y) = \mathcal{K}(y, x) = 1.0$; (2) These documents are also closely related to the other documents in the other field: $\mathcal{K}(x, y) = \mathcal{K}(y, x) = 0.9$ when x is a boundary document and y belongs to the other field; (3) These documents are not close to those in the other area in the same field: $\mathcal{K}(x, y) = \mathcal{K}(y, x) = 0.6$ when x is a boundary document and y belongs to the other area in the same field. Then the clustering $\{\text{AI, Statistics}\}$ satisfies the (α, ν) -good neighbor property only for $\alpha \geq 1/4$ or $\nu \geq 1/8$. Similarly, $\{\text{AI, ParameterEstimation, HypothesisTesting}\}$ and $\{\text{Learning, Planning, Statistics}\}$ satisfy the property only for $\alpha \geq 1/4$ or $\nu \geq 1/16$. See the text for more details, and see Section 6.1 for simulations of this example and its variants.

In this section we introduce a weaker notion of good neighborhood property and prove that our algorithm also succeeds for data satisfying this weaker property.

To motivate the property, consider a point x with the following neighborhood structure. In the neighborhood of size $n_{C(x)}$, x has a significant amount of its neighbors from other target clusters. However, in a smaller, more local neighborhood, x has most of its nearest neighbors from its target clusters $C(x)$. In practice, points close to the boundaries between different target clusters typically have such neighborhood structure; for this reason, points with such neighborhood are called boundary points.

We present an example in Figure 7. A document close to the boundary between the two fields AI and Statistics has the following neighborhood structure: out of its $n/4$ nearest neighbors, it has all its neighbors from its own field; but out of its $n/2$ nearest neighbors, it has $n/4$ neighbors outside its field. With $1/8$ fraction of such boundary points, the clustering $\{\text{AI}, \text{Statistics}\}$ satisfies the (α, ν) -good neighbor property only for $\alpha \geq 1/4$ or $\nu \geq 1/8$. This is because either we view all the boundary points as bad points in the (α, ν) -good neighborhood property which leads to $\nu \geq 1/8$, or we need $\alpha \geq 1/4$ since a boundary point has $n/4$ neighbors outside its target cluster. Similarly, $\{\text{AI}, \text{ParameterEstimation}, \text{HypothesisTesting}\}$ and $\{\text{Learning}, \text{Planning}, \text{Statistics}\}$ satisfy the property only for $\alpha \geq 1/4$ or $\nu \geq 1/16$.

For this example, either α is too large so that Theorem 1 is not applicable, or ν is too large so that the guarantee in Theorem 1 leads to constant error rate. However, it turns out that our algorithm can still successfully produce a hierarchy as in Figure 7(b) where the desired clusterings ($\{\text{Learning}, \text{Planning}, \text{Statistics}\}$, $\{\text{AI}, \text{ParameterEstimation}, \text{HypothesisTesting}\}$, $\{\text{Learning}, \text{Planning}, \text{ParameterEstimation}, \text{HypothesisTesting}\}$, and $\{\text{AI}, \text{Statistics}\}$) are prunings of the hierarchy. As we show, the reason is that each of these prunings satisfies a generalization of the good neighborhood property which takes into account the boundary points, and for which our algorithm still succeeds. Note that the standard linkage algorithms fail on this example². In the following, we first formalize this property and discuss how it relates to the properties of the similarity function described in the paper so far. We then prove that our algorithm succeeds under this property, correctly clustering all points that are not adversarially bad.

For clarity, we first relax the α -good neighborhood to the weak (α, β) -good neighborhood defined as follows.

Property 4. *A similarity function \mathcal{K} satisfies the **weak (α, β) -good neighborhood property** for the clustering problem (S, ℓ) , if for each $p \in S$, there exists $A_p \subseteq C(p)$ of size greater than $6\alpha n$ such that $p \in A_p$ and*

- *any point in A_p has at most αn neighbors outside A_p out of the $|A_p|$ nearest neighbors;*
- *for any such subset $A_p \subseteq C(p)$, at least β fraction of points in A_p have all but at most αn nearest neighbors from $C(p)$ out of their $n_{C(p)}$ nearest neighbors.*

Informally, the first condition implies that every point falls into a sufficiently large subset of its target cluster, and points in the subset are close to each other in the sense that most of their nearest neighbors are in the subset. This condition is about the local neighborhood structure of the points. It shows that each point has a local neighborhood in which points closely relate to each other. Note that the local neighborhood should be large enough so that the membership of the point is clearly established: it should have size comparable to the number of connections to points outside (αn). Here we choose a minimum size of greater than $6\alpha n$ mainly because it guarantees that our algorithm can still succeed in the worst case. The second condition

² For any fixed non-boundary point y and fixed boundary point x in the other field, the probability that y has similarity 1.0 only with x is $\frac{2}{n}(1 - \frac{2}{n})^{n/16-1} \approx \frac{2}{n}e^{-1/8}$. Since there are $n/16$ such boundary points x and $7n/8$ such non-boundary points y , when n is sufficiently large, with high probability $n/12$ non-boundary points have similarity 1.0 with one single boundary point. Then the standard linkage algorithms (in particular, single linkage, average linkage, and complete linkage) would first merge these pairs of points with similarity 1.0. From then on, no matter how they perform, any pruning of the hierarchy produced will have error higher than $1/12$.

implies that for points in these large enough subsets, a majority of them have most of their nearest neighbors from their target cluster. This condition is about more global neighborhood structure. It shows how the subsets are closely related to those in the same target cluster in the neighborhood of size equal to the target cluster size. Note that in this more global neighborhood, we do not require all points in these subsets have most nearest neighbors from their target clusters; we allow the presence of $(1 - \beta)$ fraction of points that may have a significant number of nearest neighbors outside their target clusters.

Naturally, as we can relax the α -good neighborhood property to the (α, ν) -good neighborhood property, we can relax the weak (α, β) -good neighborhood to the weak (α, β, ν) -good neighborhood as follows. Informally, it implies that the target clustering satisfies the weak (α, β) -good neighborhood property after removing a few bad points.

Property 5. *A similarity function \mathcal{K} satisfies the **weak (α, β, ν) -good neighborhood** property for the clustering problem (S, ℓ) , if there exist a subset of points B of size at most νn , and for each $p \in S \setminus B$, there exists $A_p \subseteq C(p) \setminus B$ of size greater than $6(\alpha + \nu)n$ such that $p \in A_p$ and*

- *any point in A_p has at most αn neighbors outside A_p out of the $|A_p|$ nearest neighbors;*
- *for any such subset $A_p \subseteq C_i \setminus B$, at least β fraction of points in A_p have all but at most αn nearest neighbors from $C_i \setminus B$ out of their $|C_i \setminus B|$ nearest neighbors in $S \setminus B$.*

For convenience, we call points in B bad points. If a point in $C_i \setminus B$ has all but at most αn nearest neighbors from $C_i \setminus B$ out of its $|C_i \setminus B|$ nearest neighbors in $S \setminus B$, we call it a good point. Then the second condition in the definition can be simply stated as: any A_p has at least β fraction of good points. Note that C_i can contain points that are neither bad nor good. Such points are called boundary points, since in practice such points typically lie close to the boundaries between target clusters.

As a concrete example, consider the clustering $\{\text{AI, Statistics}\}$ in Figure 7(c). It satisfies the weak (α, β, ν) -good neighborhood property with probability at least $1 - \delta$ when the number of points $n = O(\ln \frac{1}{\delta})$. To see this, first note that for a fixed point y and a fixed boundary point x in the other field, the probability that $\mathcal{K}(y, x) = 1$ is $2/n$. Since there are $n/16$ boundary point in the other field, by Hoeffding bound, the probability that y has similarity 1 with more than $n/32$ points is bounded by $\exp\{-2 \cdot n/16 \cdot (1/2)^2\} = \exp\{-n/32\}$. By union bound, with probability at least $1 - n \exp\{-n/32\}$, no point has similarity 1 with more than $n/32$ points. Then by setting A_p as the area that p falls in, we can see that the clustering satisfies the weak (α, β, ν) -good neighborhood property for $\alpha = 1/32, \beta = 7/8$ and $\nu = 0$. Note that there may also be some adversarial bad points. Then the weak (α, β, ν) -good neighborhood property is satisfied when $\alpha = 1/32, \beta = 7/8$ and ν is the fraction of bad points. See Section 6.1 for simulations of this example and its variants.

4.1 Relating Different Versions of Good Neighborhood Properties

The relations between these properties are illustrated in Figure 8. The relations between weak good neighborhood and other properties are discussed below, while the other relations in the figure follow from the facts in Section 2.1.

By setting $A_p = C_i$ for $p \in C_i$ in the definition, we can see that the weak (α, β) -good neighborhood property is a generalization of the α -good neighborhood property when each target cluster has size greater than $6\alpha n$. Formally,

Fact 5. *If the similarity function \mathcal{K} satisfies the α -good neighborhood property for the clustering problem (S, ℓ) and $\min_i |C_i| > 6\alpha n$, then \mathcal{K} also satisfies the weak (α, β) -good neighborhood property for the clustering problem (S, ℓ) for any $0 < \beta \leq 1$.*




Figure 8: Relations between various properties. The arrows represent generalization.

Proof. If \mathcal{K} satisfies the α -good neighborhood property and $\min_i |C_i| > 6\alpha n$, then we have: for any $p \in C_i$, there exists a subset $C'_i \subseteq C_i$ of size greater than $6\alpha n$, such that out of the n_{C_i} nearest neighbors, any point $x \in C_i$ has at most αn neighbors outside C_i . So \mathcal{K} satisfies both conditions of the weak (α, β) -good neighborhood property. \square

By setting $A_p = G_i$ for $p \in G_i$ in the definition, we can see that the weak (α, β, ν) -good neighborhood property generalizes the (α, ν) -good neighborhood property when each target cluster has size greater than $7(\alpha + \nu)n$. Also, by setting $\nu = 0$, we can see that the weak (α, β) -good neighborhood property is equivalent to the weak $(\alpha, \beta, 0)$ -good neighborhood.

Fact 6. *If the similarity function \mathcal{K} satisfies the (α, ν) -good neighborhood property for the clustering problem (S, ℓ) and $\min_i |C_i| > 7(\alpha + \nu)n$, then \mathcal{K} also satisfies the weak (α, β, ν) -good neighborhood property for the clustering problem (S, ℓ) for any $0 < \beta \leq 1$.*

Proof. If \mathcal{K} satisfies the (α, ν) -good neighborhood property and $\min_i |C_i| > 7(\alpha + \nu)n$, then we have: for any $p \in G_i = C_i \setminus B$, there exists a subset $G'_i \subseteq G_i$ of size greater than $6(\alpha + \nu)n$, such that out of the $|G_i|$ nearest neighbors in $S \setminus B$, any good point $x \in G_i$ has at most αn neighbors outside G_i . So \mathcal{K} satisfies both conditions of the weak (α, β, ν) -good neighborhood property. \square

Fact 7. *If the similarity function \mathcal{K} satisfies the weak (α, β) -good neighborhood property for the clustering problem (S, ℓ) , then \mathcal{K} also satisfies the weak $(\alpha, \beta, 0)$ -good neighborhood property for the clustering problem (S, ℓ) .*

Proof. By setting $\nu = 0$ in the definition of the weak (α, β, ν) -good neighborhood property, we can see that it is the same as the weak (α, β) -good neighborhood property. \square

4.2 Correctness under the Weak Good Neighborhood Property

Now we prove that our algorithm also succeeds under the weak (α, β, ν) -good neighborhood property when $\beta \geq 7/8$. Formally,

Theorem 2. *Let \mathcal{K} be a symmetric similarity function satisfying the weak (α, β, ν) -good neighborhood property for the clustering problem (S, ℓ) with $\beta \geq 7/8$. Then Algorithm 1 outputs a hierarchy such that a pruning of the hierarchy is ν -close to the target clustering in time $O(n^{\omega+1})$, where $O(n^\omega)$ is the state of the art for matrix multiplication.*




Figure 9: Illustration of a fully formed blob C_u : for any point $p \in C_u \setminus B$, $A_p \subseteq C_u$. Then we can show that sets in $\{A_p : p \in C_u \setminus B\}$ are laminar, that is, for any $p, q \in C_u \setminus B$, either $A_p \cap A_q = \emptyset$ or $A_p \subseteq A_q$ or $A_q \subseteq A_p$. For example, in the figure we have $A_{p_5} \subseteq A_{p_4}$.

The correctness of the theorem follows from Lemma 5 stated and proved below and the running time follows from Lemma 2. The intuition is as follows. First, by similar arguments as for good neighborhood property, each point p in $S \setminus B$ will only be merged with other points in A_p at $t \leq |A_p|$, and all points in A_p will belong to one blob at $t = |A_p|$ (Lemma 3), since in the local neighborhood of size $|A_p|$, the point has most of its nearest neighbor from A_p . Then, we need to show that such blobs will be correctly merged. The key point is to show that even in the presence of boundary points, the majority of points in such blobs are good points (Lemma 4). Then the median test can successfully distinguish blobs containing good points from different target clusters, and our algorithm can correctly merge blobs from the same target clusters together.

To formally prove the correctness, we begin with Lemma 3. The proof is similar to that for Lemma 1, replacing C_i with A_p .

Lemma 3. *The following claims are true in Algorithm 1:*

- (1) *For any point $p \in S \setminus B$ and t such that $t \leq |A_p|$, any blob in \mathcal{C}'_t containing points from A_p will not contain points in $(S \setminus A_p) \setminus B$.*
- (2) *For any point $p \in S \setminus B$ and $t = |A_p|$, all points in A_p belong to one blob in \mathcal{C}'_t .*

Lemma 3 states that for any $p \in S \setminus B$, we will form A_p before merging them with points outside. Then we only need to make sure that these A_p formed will be correctly merged. More precisely, we need to consider the blobs that are “fully formed” in the following sense:

Definition 1. *A blob $C_u \in \mathcal{C}'_t$ in Algorithm 1 is said to be fully formed if for any point $p \in C_u \setminus B$, $A_p \subseteq C_u$.*

To show that fully formed blobs are correctly merged, the key point is to show that the majority of points in such blobs are good points, and thus the median test in the algorithm can successfully distinguish blobs containing good points from different target clusters. This key point is in fact a consequence of Lemma 3:

Lemma 4. *For any fully formed blob $C_u \in \mathcal{C}'_t$ in Algorithm 1, at least β fraction of points in $C_u \setminus B$ are good points.*

Proof. It suffices to show that there exist a set of points $P \subseteq C_u \setminus B$, such that $\{A_p : p \in P\}$ is a partition of $C_u \setminus B$. Clearly $C_u \setminus B = \cup_{p \in C_u \setminus B} A_p$. So we only need to show that sets in $\{A_p : p \in C_u \setminus B\}$ are

laminar, that is, for any $p, q \in C_u \setminus B$, either $A_p \cap A_q = \emptyset$ or $A_p \subseteq A_q$ or $A_q \subseteq A_p$. See Figure 9 for an illustration.

Assume for contradiction that there exist A_p and A_q such that $A_p \setminus A_q \neq \emptyset$, $A_q \setminus A_p \neq \emptyset$ and $A_p \cap A_q \neq \emptyset$. Without loss of generality, suppose $|A_p| \leq |A_q|$. Then by the second claim in Lemma 3, when $t = |A_p|$, all points in A_p belong to one blob in \mathcal{C}'_t . In other words, this blob contains $A_p \cap A_q$ and $A_p \setminus A_q$. So for $t \leq |A_q|$, the blob contains points in A_q and also points in $S \setminus B \setminus A_q$, which contradicts the first claim in Lemma 3. \square

We are now ready to prove the following lemma that implies Theorem 2.

Lemma 5. *The following claims are true in Algorithm 1:*

- (1) *For any C_i such that $t \leq |C_i|$, any blob in \mathcal{C}'_t containing points in $C_i \setminus B$ will not contain points in $(S \setminus C_i) \setminus B$.*
- (2) *For any C_i such that $t = |C_i|$, all points in $C_i \setminus B$ belong to one blob in \mathcal{C}'_t .*

Proof. Before proving the claims, we first show that the graph F_t constructed in Step 2 has the following useful properties³. Recall that F_t is constructed on points in S by connecting any two points that share at least $t - 2(\alpha + \nu)n$ points in common out of their t nearest neighbors. For any C_i such that $t \leq |C_i|$, we have:

- (a) If x is a good point in C_i and y is a good point outside C_i , then x and y are not connected in F_t .

To see this, first note that by Fact 3, x has at most $(\alpha + \nu)n$ neighbors outside C_i out of the t nearest neighbors. Suppose y is a good point from C_j . If $n_{C_j} \geq t$, then y has at most $(\alpha + \nu)n$ neighbors in C_i ; if $n_{C_j} < t$, y has at most $(\alpha + \nu)n + t - n_{C_j}$ neighbors in C_i . In both cases, y has at most $(\alpha + \nu)n + \max(0, t - n_{C_j}) < t - 5(\alpha + \nu)n$ neighbors in C_i , since $n_{C_j} > 6(\alpha + \nu)n$ and $t > 6(\alpha + \nu)n$. Then x and y have at most $t - 4(\alpha + \nu)n$ common neighbors, so they are not connected in F_t .

- (b) If x is a good point in C_i , y is a good point outside C_i , and z is a bad point, then z cannot be connected to both x and y in F_t .

To prove this, we will show that if z is connected to x , then z cannot be connected to y . First, by the same argument as above, out of the t nearest neighbors, y has less than $t - 5(\alpha + \nu)n$ neighbors in C_i . Second, by Fact 3, x has at most $(\alpha + \nu)n$ neighbors outside C_i . If z has less than $t - 3(\alpha + \nu)n$ neighbors in C_i , then z and x share less than $t - 3(\alpha + \nu)n + (\alpha + \nu)n = t - 2(\alpha + \nu)n$ neighbors and will not be connected. So z must have at least $t - 3(\alpha + \nu)n$ neighbors in C_i , and thus cannot have more than $3(\alpha + \nu)n$ neighbors outside C_i . The two statements show that y and z share less than $t - 5(\alpha + \nu)n$ neighbors in C_i , and at most $3(\alpha + \nu)n$ neighbors outside C_i . So they share less than $t - 2(\alpha + \nu)n + 3(\alpha + \nu)n = t - 2(\alpha + \nu)n$ neighbors and thus are not connected in F_t .

Now we prove Claim (1) in the lemma by induction on t . The claim is clearly true initially. Assume for induction that the claim is true for the threshold $t - 1$, that is, for any C_i such that $t - 1 \leq |C_i|$, any blob in \mathcal{C}'_{t-1} containing points in $C_i \setminus B$ will not contain points in $(S \setminus C_i) \setminus B$. We now prove that the graph H_t constructed in Step 3 has the following properties, which can be used to show that the claim is still true for the threshold t .

³The argument is similar to that in Lemma 1.

- If $C_u \in \mathcal{C}'_{t-1}$ contains points from $C_i \setminus B$ and $C_v \in \mathcal{C}'_{t-1}$ contains points from $(S \setminus C_i) \setminus B$, then they cannot be connected in H_t .

Suppose one of them (say C_u) is not fully formed, that is, there is a point $p \in C_u \setminus B$ such that $A_p \not\subseteq C_u$. Then by Lemma 3, the algorithm will not merge C_u with C_v at this threshold. More precisely, since not all points in A_p belong to C_u , we have $t - 1 < |A_p|$ by Claim (2) in Lemma 3. Then by Claim (1) in Lemma 3, since C_v contains points in $(S \setminus A_p) \setminus B$, C_u and C_v will not be merged in \mathcal{C}'_t . Then they cannot be connected in H_t .

So we only need to consider the other case when C_u and C_v are fully formed blobs. By Lemma 4, the majority of points in the two blobs are good points. The good points from different target clusters have few common neighbors in F_t , then by the median test in our algorithm, the two blobs will not be connected in H_t . Formally, we can find two good points $x^* \in C_u, y^* \in C_v$ that satisfy the following two statements.

- $S_t(x^*, y^*) \geq \text{median}_{x \in C_u, y \in C_v} S_t(x, y)$.

By Lemma 4, at least $\beta \geq 7/8$ fraction of points in $C_u \setminus B$ are good points. The fraction of good points in C_u is at least

$$\frac{\beta|C_u \setminus B|}{|C_u \setminus B| + |B|} \geq \frac{7/8 \times 6(\alpha + \nu)n}{6(\alpha + \nu)n + \nu n} \geq \frac{3}{4}$$

since $|C_u \setminus B| \geq 6(\alpha + \nu)n$ and $|B| \leq \nu n$. Similarly, at least $\frac{3}{4}$ fraction of points in C_v are good points. Then among all the pairs (x, y) such that $x \in C_u, y \in C_v$, at least $\frac{3}{4} \times \frac{3}{4} > \frac{1}{2}$ fraction are pairs of good points. So there exist good points $x^* \in C_u, y^* \in C_v$ such that $S_t(x^*, y^*) \geq \text{median}_{x \in C_u, y \in C_v} S_t(x, y)$.

- $S_t(x^*, y^*) \leq (|C_u| + |C_v|)/4$.

The fraction of good points in $C_u \cup C_v$ is at least $\frac{3}{4}$. Since in F_t , good points in C_u are not connected to good points in C_v , we have $S_t(x^*, y^*) \leq (|C_u| + |C_v|)/4$.

Combining the two statements, we have $\text{median}_{x \in C_u, y \in C_v} S_t(x, y) \leq (|C_u| + |C_v|)/4$ and thus C_u and C_v are not connected in H_t .

- If in \mathcal{C}'_{t-1} , C_u contains points from $C_i \setminus B$, C_v contains points from $(S \setminus C_i) \setminus B$, and C_w contains only bad points, then C_w cannot be connected to both C_u and C_v .

By the same argument as above, we only need to consider the case when C_u and C_v are fully formed blobs. To prove the claim in this case, assume for contradiction that C_w is connected to both C_u and C_v . First, note the following fact about C_w . Since any non-singleton blob must be formed in Step 4 in the algorithm and contain at least $4(\alpha + \nu)n$ points and thus cannot contain only bad points, C_w must be a singleton blob, containing only a bad point z .

Next, we show that if $C_w = \{z\}$ are connected to C_u in H_t , then z must be connected to at least one good point in C_u in F_t . We have $\text{median}_{x \in C_u} S_t(x, z) > \frac{|C_u| + |C_w|}{4}$, which means z is connected to more than $\frac{|C_u|}{4}$ points in C_u in F_t . By the same argument as above, at least $3/4$ fraction of points in C_u are good points, then z must be connected to at least one good point in C_u .

Similarly, if C_w is connected to C_v in H_t , then z must be connected to at least one good point in C_v in F_t . But this contradicts Property (b) of F_t , so C_w cannot be connected to both C_u and C_v in H_t .

By the properties of H_t , no connected component contains both points in $C_i \setminus B$ and points in $(S \setminus C_i) \setminus B$. So Claim (1) is still true for the threshold t . By induction, it is true for all thresholds.

Finally, we prove Claim (2). By Lemma 3, when $t = |C_i|$, for any point $p \in C_i \setminus B$, A_p belong to the same blob. So all points in $C_i \setminus B$ are in sufficiently large blobs. We will show that any two of these blobs C_u, C_v are connected in H_t , and thus will be merged into one blob. By Lemma 4, we know that more than $3/4$ fraction of points in C_u (C_v respectively) are good points, and thus there exist good points $x^* \in C_u, y^* \in C_v$ such that $S_t(x^*, y^*) \leq \text{median}_{x \in C_u, y \in C_v} S_t(x, y)$. By Claim (1), all good points in C_u and C_v are from C_i , so they share at least $t - 2(\alpha + \nu)n$ neighbors when $t = |C_i|$, and thus are connected in F_t . Then $S_t(x^*, y^*)$ is at least the number of good points in $C_u \cup C_v$, which is at least $3(|C_u| + |C_v|)/4$. Then $\text{median}_{x \in C_u, y \in C_v} S_t(x, y) \geq S_t(x^*, y^*) > (|C_u| + |C_v|)/4$. Therefore, all blobs containing points from $C_i \setminus B$ are connected in H_t and thus merged into a blob. \square

5 The Inductive Setting

Many clustering applications have recently faced an explosion of data, such as in astrophysics and biology. For large data sets, it is often resource and time intensive to run an algorithm over the entire data set. It is thus increasingly important to develop algorithms that can remove the dependence on the actual size of the data and still perform reasonably well.

In this section we consider an inductive model that formalizes this problem. In this model, the given data is merely a small random subset of points from a much larger data set. The algorithm outputs a hierarchy over the sample, which also implicitly represents a hierarchy over the data set. In the following we describe the inductive version of our algorithm and prove that when the data satisfies good neighborhood properties, the algorithm achieves small error on the entire data set, requiring only a small random sample whose size is independent of that of the entire data set.

5.1 Formal Definition

First we describe the formal definition of the inductive model. In this setting, the given data S is merely a small random subset of points from a much larger abstract instance space X . For simplicity, we assume that X is finite and that the underlying distribution is uniform over X . Let $N = |X|$ denote the size of the entire instance space, and let $n = |S|$ denote the size of the sample.

Our goal is to design an algorithm that based on the sample produces a hierarchy of small error with respect to the whole distribution. Formally, we assume that each node u in the hierarchy derived over the sample induces a cluster (a subset of X). For convenience, u is also used to denote the blob of sampled points it represents. The cluster u induces over X is implicitly represented as a function $f_u : X \rightarrow \{0, 1\}$, that is, for each $x \in X$, $f_u(x) = 1$ if x is a point in the cluster and 0 otherwise. We say that the hierarchy has error at most ϵ if it has a pruning f_{u_1}, \dots, f_{u_k} of error at most ϵ .

5.2 Inductive Robust Median Neighborhood Linkage

The inductive version of our algorithm is described in Algorithm 2. To analyze the algorithm, we first present the following lemmas showing that, when the data satisfies the good neighborhood property, a sample of sufficiently large size also satisfies the weak good neighborhood property.

Lemma 6. *Let \mathcal{K} be a symmetric similarity function satisfying the (α, ν) -good neighborhood for the clustering problem (X, ℓ) . Consider any fixed $x \in X \setminus B$. If the size of the sample satisfies $n = \Theta(\frac{1}{\alpha} \ln \frac{1}{\delta})$, then with probability at least $1 - \delta$, x has at most $2\alpha n$ neighbors outside $(C(x) \setminus B) \cap S$ out of the $|C(x) \setminus B|$ nearest neighbors in $S \setminus B$.*

Algorithm 2 Inductive Robust Median Neighborhood Linkage

Input: similarity function \mathcal{K} , $n \in \mathbf{Z}^+$, parameters $\alpha > 0, \nu > 0$.

▷ Get a hierarchy on the sample

Sample i.i.d. examples $S = \{x_1, \dots, x_n\}$ uniformly at random from X .

Run Algorithm 1 with parameters $(2\alpha, 2\nu)$ on S and obtain a hierarchy T .

▷ Get the implicit hierarchy over X

for any $x \in X$ **do**

 Let $N_S(x)$ denote the $6(\alpha + \nu)n$ nearest neighbors of x in S .

 Initialize $u = \text{root}(T)$ and $f_u(x) = 1$.

while u is not a leaf **do**

 Let w be the child of u that contains the most points in $N_S(x)$.

 Set $u = w$ and $f_u(x) = 1$.

end while

end for

Output: Hierarchy T and $\{f_u, u \in T\}$.

Proof. Suppose $x \in G_i$. Let $NN(x)$ denote its $|G_i|$ nearest neighbors in X . By assumption we have that $|NN(x) \setminus G_i| \leq \alpha N$ and $|G_i \setminus NN(x)| \leq \alpha N$. Then by Chernoff bounds, with probability at least $1 - \delta$ at most $2\alpha n$ points in our sample are in $NN(x) \setminus G_i$ and at most $2\alpha n$ points in our sample are in $G_i \setminus NN(x)$.

We now argue that at most $2\alpha n$ of the $|G_i \cap S|$ nearest neighbors of x in $S \setminus B$ can be outside G_i . Let n_1 be the number of points in $(NN(x) \setminus G_i) \cap S$, n_2 be the number of points in $(G_i \setminus NN(x)) \cap S$, and n_3 be the number of points in $(G_i \cap NN(x)) \cap S$. Then $|G_i \cap S| = n_2 + n_3$ and we know that $n_1 \leq 2\alpha n$, $n_2 \leq 2\alpha n$. We consider the following two cases.

- $n_1 \geq n_2$. Then $n_1 + n_3 \geq n_2 + n_3 = |G_i \cap S|$. This implies that the $|G_i \cap S|$ nearest neighbors of x in the sample all lie inside $NN(x)$, since by definition all points inside $NN(x)$ are closer to x than any point outside $NN(x)$. But we are given that at most $n_1 \leq 2\alpha n$ of them can be outside G_i . Thus, we get that at most $2\alpha n$ of the $|G_i \cap S|$ nearest neighbors of x are not from G_i .
- $n_1 < n_2$. This implies that the $|G_i \cap S|$ nearest neighbors of x in the sample include *all* the points in $NN(x)$ in the sample, and possibly some others too. But this implies in particular that it includes all the n_3 points in $G_i \cap NN(x)$ in the sample. So, it can include at most $|G_i \cap S| - n_3 = n_2 \leq 2\alpha n$ points not in $G_i \cap NN(x)$. Even if all those are not in G_i , the $|G_i \cap S|$ nearest neighbors of x still include at most $2\alpha n$ points not from G_i .

In both cases, at most $2\alpha n$ of the $|G_i \cap S|$ nearest neighbors of x in $S \setminus B$ can be outside G_i . □

Lemma 7. Let \mathcal{K} be a symmetric similarity function satisfying the (α, ν) -good neighborhood for the clustering problem (X, ℓ) . If the size of the sample satisfies $n = \Theta\left(\frac{1}{\min(\alpha, \nu)} \ln \frac{1}{\delta \min(\alpha, \nu)}\right)$, then with probability at least $1 - \delta$, \mathcal{K} satisfies the $(2\alpha, 2\nu)$ -good neighborhood with respect to the clustering induced over the sample S .

Proof. First, by Chernoff bounds, when $n \geq \frac{3}{\nu} \ln \frac{2}{\delta}$, we have that with probability at least $1 - \delta/2$, at most $2\nu n$ bad points fall into the sample. Next, by Lemma 6 and union bound, when $n = \Theta\left(\frac{1}{\alpha} \ln \frac{n}{\delta}\right)$ we have that

with probability at least $1 - \delta/2$, for any C_i , any $x \in G_i \cap S$, x has at most $2\alpha n$ points outside $G_i \cap S$ out of its $|G_i \cap S|$ nearest neighbors in $(X \setminus B) \cap S$. Therefore, if $n = \Theta\left(\frac{1}{\min(\alpha, \nu)} \ln \frac{n}{\delta}\right)$, then with probability at least $1 - \delta$, the similarity function satisfies the $(2\alpha, 2\nu)$ -good neighborhood property with respect to the clustering induced over the sample S .

It now suffices to show n is large enough so that $n = \Theta\left(\frac{1}{\min(\alpha, \nu)} \ln \frac{n}{\delta}\right)$. To see this, let $\eta = \min(\alpha, \nu)$. Since $\ln n \leq tn - \ln t - 1$ for any $t, n > 0$, we have

$$\frac{c}{\eta} \ln n \leq \frac{c}{\eta} \left(\frac{\eta}{2c} n + \ln \frac{2c}{\eta} - 1 \right) = \frac{n}{2} + \frac{c}{\eta} \ln \frac{2c}{e \cdot \eta}$$

for any constant $c > 0$. Then $n = \Theta\left(\frac{1}{\eta} \ln \frac{1}{\eta}\right)$ implies $n = \Theta\left(\frac{1}{\eta} \ln n\right)$, and $n = \Theta\left(\frac{1}{\eta} \ln \frac{1}{\delta \cdot \eta}\right)$ implies $n = \Theta\left(\frac{1}{\eta} \ln \frac{n}{\delta}\right)$. \square

Theorem 3. *Let \mathcal{K} be a symmetric similarity function satisfying the (α, ν) -good neighborhood for the clustering problem (X, ℓ) . As long as the smallest target cluster has size greater than $12(\nu + \alpha)N$, then Algorithm 2 with parameters $n = \Theta\left(\frac{1}{\min(\alpha, \nu)} \ln \frac{1}{\delta \cdot \min(\alpha, \nu)}\right)$ produces a hierarchy with a pruning that is $(\nu + \delta)$ -close to the target clustering with probability $1 - \delta$.*

Proof. Note that by Lemma 7, with probability at least $1 - \delta/4$, we have that \mathcal{K} satisfies the $(2\alpha, 2\nu)$ -good neighborhood with respect to the clustering induced over the sample. Moreover, by Chernoff bounds, with probability at least $1 - \delta/4$, each G_i has at least $6(\nu + \alpha)n$ points in the sample. Then by Theorem 1, Algorithm 1 outputs a hierarchy T on the sample S with a pruning that assigns all good points correctly. Denote this pruning as $\{u_1, \dots, u_k\}$ such that $u_i \setminus B = (C_i \cap S) \setminus B$.

Now we want to show that f_{u_1}, \dots, f_{u_k} have error at most $\nu + \delta$ with probability at least $1 - \delta/2$. For convenience, let $u(x)$ be a shorthand of $u_{\ell(x)}$. Then it is sufficient to show that with probability at least $1 - \delta/2$, a $(1 - \delta)$ fraction of points $x \in X \setminus B$ have $f_{u(x)}(x) = 1$. Fix C_i and a point $x \in C_i \setminus B$. By Lemma 6, with probability at least $1 - \delta^2/2$, out of the $|G_i \cap S|$ nearest neighbors of x in $S \setminus B$, at most $2\alpha n$ can be outside G_i . Recall that Algorithm 2 checks $N_S(x)$, the $6(\alpha + \nu)n$ nearest neighbors of x in S . Then out of $N_S(x)$, at most $2(\alpha + \nu)n$ points are outside $G_i \cap S$. By Lemma 1, u_i contains $G_i \cap S$, so u_i must contain at least $4(\alpha + \nu)n$ points in $N_S(x)$. Consequently, any ancestor w of u_i , including u_i , has more points in $N_S(x)$ than any other sibling of w . Then we must have $f_w(x) = 1$ for any ancestor w of u_i . In particular, $f_{u_i}(x) = 1$. So, for any point $x \in X \setminus B$, with probability at least $1 - \delta^2/2$ over the draw of the random sample, $f_{u(x)}(x) = 1$. By Markov inequality, with probability at least $1 - \delta/2$, a $(1 - \delta)$ fraction of points $x \in X \setminus B$ have $f_{u(x)}(x) = 1$. More precisely, let U_x denote the uniform distribution over $X \setminus B$, and let U_S denote the distribution of the sample S . Let $I(x, S)$ denote the event that $f_{u(x)}(x) \neq 1$. Then we have

$$\mathbf{E}_{x \sim U_x, S \sim U_S}[I(x, S)] = \mathbf{E}_{S \sim U_S} \left[\mathbf{E}_{x \sim U_x}[I(x, S)|S] \right] \leq \delta^2/2.$$

Then by Markov inequality, we have

$$\Pr_{S \sim U_S} \left[\mathbf{E}_{x \sim U_x}[I(x, S)|S] \geq \delta \right] \leq \delta/2$$

which means that with probability at least $1 - \delta/2$ over the draw of the random sample S , a $(1 - \delta)$ fraction of points $x \in X \setminus B$ have $f_{u(x)}(x) = 1$. \square

Similarly, Algorithm 2 also succeeds for weak good neighborhood. By similar arguments as those in Lemma 6 and 7, we can prove that the similarity \mathcal{K} satisfies weak good neighborhood property over a sufficiently large sample (Lemma 8), which then leads to the final guarantee Theorem 4. For clarity, the proofs are provided in Appendix B.

Lemma 8. *Let \mathcal{K} be a symmetric similarity function satisfying the weak (α, β, ν) -good neighborhood for the clustering problem (X, ℓ) . Furthermore, it satisfies that for any $p \in X \setminus B$, $|A_p| > 24(\alpha + \nu)N$. If the size of the sample satisfies $n = \Theta\left(\frac{1}{\min(\alpha, \nu)} \ln \frac{1}{\delta \min(\alpha, \nu)}\right)$, then with probability at least $1 - \delta$, \mathcal{K} satisfies the $(2\alpha, \frac{15}{16}\beta, 2\nu)$ -good neighborhood with respect to the clustering induced over the sample S .*

Theorem 4. *Let \mathcal{K} be a symmetric similarity function satisfying the weak (α, β, ν) -good neighborhood for the clustering problem (X, ℓ) with $\beta \geq \frac{14}{15}$. Furthermore, it satisfies that for any $p \in X \setminus B$, $|A_p| > 24(\alpha + \nu)N$. Then Algorithm 2 with parameters $n = \Theta\left(\frac{1}{\min(\alpha, \nu)} \ln \frac{1}{\delta \min(\alpha, \nu)}\right)$ produces a hierarchy with a pruning that is $(\nu + \delta)$ -close to the target clustering with probability $1 - \delta$.*

6 Experiments

In this section, we compare our algorithm (called RMNL for convenience) with popular hierarchical clustering algorithms, including standard linkage algorithms (Sneath and Sokal, 1973; King, 1967; Everitt et al., 2011), (Generalized) Wishart’s Method (Wishart, 1969; Chaudhuri and Dasgupta, 2010), Ward’s method (Ward, 1963), CURE (Guha et al., 1998), and EigenCluster (Cheng et al., 2006).

To evaluate the performance of the algorithms, we use the model discussed in Section 2. Given a hierarchy output by an algorithm, we generate all possible prunings of size k , where k is the number of clusters in the target clustering. Then we compute the *Classification Error* of each pruning with respect to the target clustering, and report the best error. The Classification Error of a computed clustering h with respect to the target clustering ℓ is the probability that a point chosen at random from the data is labeled incorrectly⁴. Formally,

$$\text{err}(h) = \min_{\sigma \in \mathcal{S}_k} \left[\Pr_{x \in S} [\sigma(h(x)) \neq \ell(x)] \right]$$

where \mathcal{S}_k is the set of all permutations on $\{1, \dots, k\}$.

Data sets To emphasize the effect of noise on different algorithms, we perform controlled experiments on a synthetic data set AStat. This data set contains 512 points. It is an instance of the example discussed in Section 4 and is described in Figure 7. We further consider the following real-world data sets from UCI Repository (Bache and Lichman, 2013): Wine, Iris, BCW (Breast Cancer Wisconsin), BCWD (Breast Cancer Wisconsin Diagnostic), Spambase, and Mushroom. We also consider the MNIST data set (LeCun et al., 1998) and use two subsets of the test set for our experiments: Digits0123 that contains the examples of the digits 0, 1, 2, 3, and Digits4567 that contains the examples of the digits 4, 5, 6, 7.

We additionally consider the 10 data sets (PFAM1 to PFAM10) from (Voevodski et al., 2012), which are created by randomly choosing 8 families (of size between 1000 and 10000) from the biology database Pfam (Punta et al., 2012), version 24.0, October 2009. The similarities for the PFAM data sets are generated by biological sequence alignment software BLAST (Altschul et al., 1990). BLAST performs one versus

⁴To compute this error for a computed clustering in polynomial time, we first find its best match to the target clustering using the Hungarian Method (Kuhn, 1955) for min-cost bipartite matching in time $O(n^3)$, and then calculate the error as the fraction of points misclassified in matched clusters.

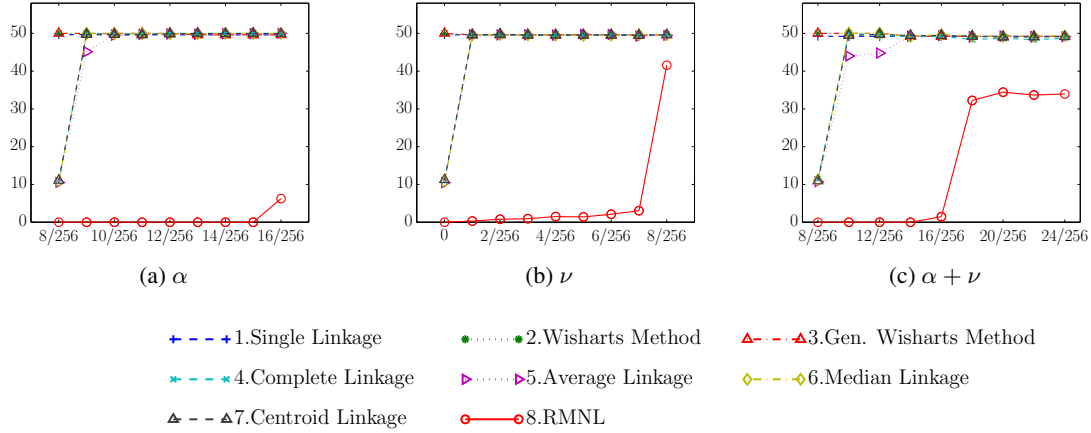


Figure 10: Classification Error on the synthetic data AISTat. The y -axis represents the % error. (a) Fix $\nu = 0$ and vary α from $1/32$ to $1/16$. The x -axis represents the value of α ; (b) Fix $\alpha = 1/32$ and vary ν from 0 to $1/32$. The x -axis represents the value of ν ; (c) Vary α from $1/32$ to $1/16$, and vary ν from 0 to $1/32$. The x -axis represents the value of $\alpha + \nu$. Note that the instance no longer satisfies the weak good neighborhood property when $\alpha + \nu \geq 1/24 \approx 11/256$.

all queries by aligning a queried sequence to sequences in the data set, and produces a score for each alignment. The score is a measure of the alignment quality and thus can be used as similarity. However, BLAST does not consider alignments with some of the sequences, in which case we assign similarities 0 to the corresponding sequences and exclude them from the neighbors of the queried sequence.

Among these data sets, the smaller ones (Wine, Iris, BCW, BCWD, and Ionosphere) are used in the transductive setting, and the larger ones (Spambase, Mushroom, Digits0123, Digits4567, and PFAM1 to PFAM10) are used in the inductive setting.

6.1 Synthetic Data

Here we compare the performance of the algorithms on the synthetic data AISTat. Recall that the clustering $\{\text{AI, Statistics}\}$ satisfies the weak (α, β, ν) -good neighborhood property for $\alpha = 1/32, \beta = 7/8, \nu = 0$ with high probability (See Figure 7 in Section 4). We conduct three sets of experiments, where we vary the values of α and ν by modifying the similarities between the points.

- (a) For each point x , we choose $\Delta\alpha n$ points y from the other field and set the similarities $\mathcal{K}(x, y) = \mathcal{K}(y, x) = 1$, so that the value of α is increased to $1/32 + \Delta\alpha$. By varying $\Delta\alpha$, we control $\alpha = 1/32 + i/256$ for $i = 0, \dots, 8$ and run the clustering algorithms on the modified data.
- (b) We randomly choose νn points x , and then set the similarity between x and any other point to be 1 minus the original similarity. This introduces νn bad points. We thus control $\nu = i/256$ for $i = 0, \dots, 8$.
- (c) We perform the above two modifications simultaneously, that is, we control $\alpha = 1/32 + i/256$ and $\nu = i/256$ for $i = 0, \dots, 8$.

Note that the instance no longer satisfies the weak good neighborhood property when $\alpha + \nu \geq 1/24$. This is because the weak good neighborhood requires that each point $p \notin B$ falls into a subset A_p of size greater

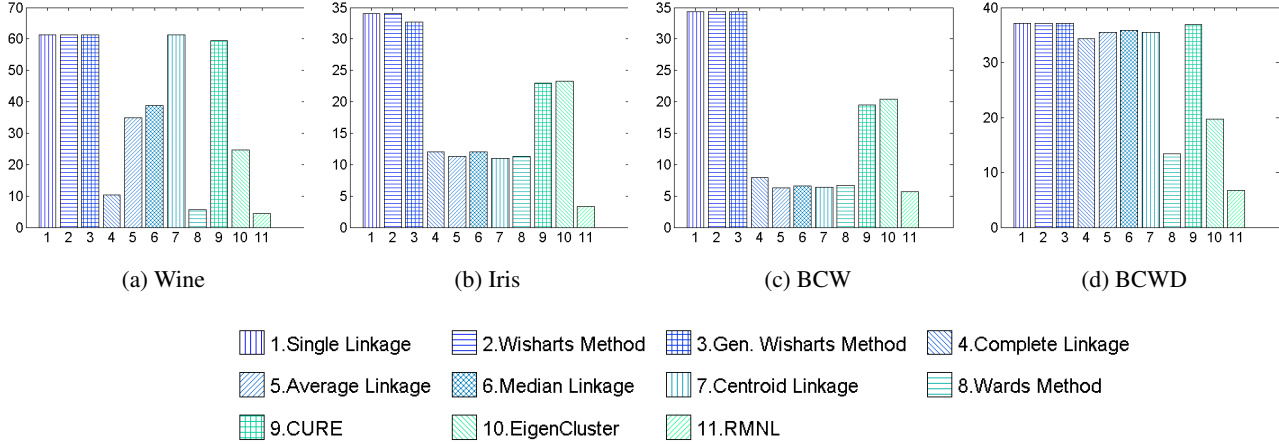


Figure 11: Classification Error in the transductive setting. The y -axis represents the % error.

than $6(\alpha + \nu)n$ with desired properties (see Property 5), and the largest such subsets in this instance have size $n/4$.

Figure 10 shows the results of these experiments, averaged over 10 runs. When $\alpha + \nu < 1/24$, the instance satisfies the weak good neighborhood property and our algorithm has error at most ν . Moreover, even if the instance does not satisfy the weak good neighborhood property when $\alpha + \nu \geq 1/24$, our algorithm still reports lower error. All the other algorithms have higher error than our algorithm and fail rapidly as $\alpha + \nu$ increases. This demonstrates that in cases modeled by the properties we propose, our algorithm will be successful while the traditional agglomerative algorithms fail.

6.2 Real-World Data

In this section, we compare the performance of our algorithm with the other algorithms on real-world data sets and show that our algorithm consistently outperforms the others. For reporting results, we follow the methodology in (Guha et al., 1998): for all algorithms accepting input parameters (including (Generalized) Wisharts' Method, CURE, and RMNL), the experiments are repeated on the same data over a range of input parameter values, and the best results are considered.

6.2.1 Transductive Setting

Here we compare the performance of the algorithms in the transductive setting where the algorithms use all the points in the data set. Figure 11 shows that our algorithm consistently achieves lowest or close to lowest errors on all the data sets. Ward's Method is the best among the other algorithms, but still shows larger errors than our algorithms. All the other algorithms generally show worse performance, and report significantly higher errors on some of the data sets. The comparison shows the robustness of our algorithm to the noise in the real world data sets.

To further evaluate the robustness of the algorithms, in the following we show their performance when different types of noise are added to the data. Since our algorithm requires additional parameters to characterize noise, we also discuss their robustness to parameter tuning.

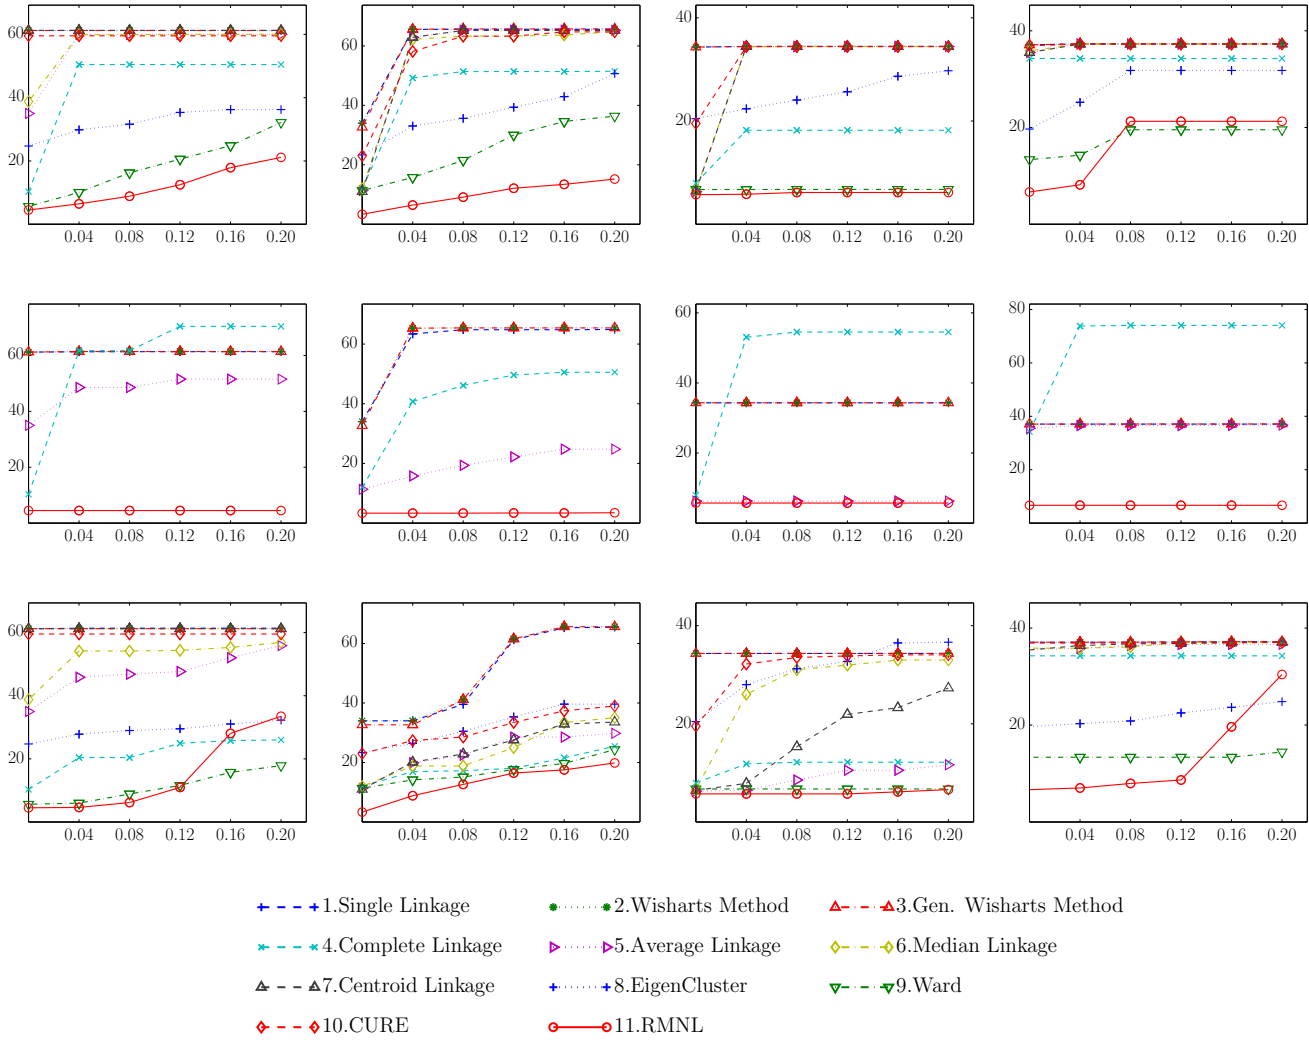


Figure 12: Classification Error in the presence of noise. Rows: corruption noise added to the attributes, corruption noise added to the similarities, Gaussian noise added to the attributes. Columns: Wine, Iris, BCW, BCWD. In each subfigure, the x -axis represents the noise level, and the y -axis represents the % error.

Robustness to Noise Here we present the performance of the algorithms when Gaussian noise or corruption noise is added and the level of noise is increased monotonically. The Gaussian noise models additive perturbations to the data entries, and is probably the most common type of noise studied. The corruption noise models data corruption or missing values, and is also frequently studied in machine learning and coding theory (Blum et al., 2007; Feldman et al., 2008; Wigderson and Yehudayoff, 2012; Moitra and Saks, 2013). The experiments on different types of noise then evaluate the robustness of the algorithms to noise caused by different reasons in real world scenarios. Note that the instance is not in a metric space after adding noise to the similarities, so in this case, we only evaluate algorithms that can be run on non-metric instances.

We consider three types of noise: corruption noise to the attributes, corruption noise to the similarities, and Gaussian noise added to the attributes. The first type of noise is generated as follows: normalize the

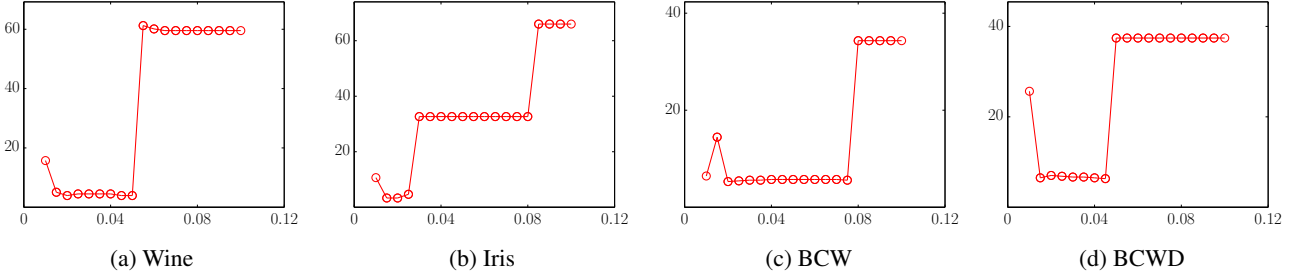


Figure 13: Classification Error using different values of parameter $(\alpha + \nu)$. The x -axis represents the value of $(\alpha + \nu)$, and the y -axis represents the % error.

entries in the data matrix to $[0, 1]$; randomly pick p fraction of the entries; replace each sampled entry with a random value independently generated from $N(0, 1)$, where p is the parameter indicating the level of noise. The second type of noise is generated using the same approach, but is added to the similarity matrix. The third type of noise is generated as follows: normalize the entries in the data matrix to $[0, 1]$; add a random value independently generated from $N(0, p^2)$ to each entry, where p is the parameter indicating the level of noise.

Figure 12 shows the results of different algorithms in the presence of noise, averaged over 30 runs. The rows correspond to different types of noise added, and the columns correspond to different data sets. The first row shows the results when corruption noise is added to the attributes. Our algorithm shows robustness to such type of noise: its error rates remain the best or close to the best up to noise level 0.2 on all data sets. EigenCluster and Ward’s method also show robustness, but their error rates are generally higher than those of our algorithm. The other algorithms report high errors even when the noise level is as low as 0.04.

The second row shows the results when corruption noise is added to the similarities. We observe that the errors of our algorithm remain nearly the same up to noise level 0.2 over all the data sets, while the other algorithms report higher errors. Some algorithms (such as Complete Linkage on Wine) show comparable performance to our algorithms when there is no noise, but their errors generally increase rapidly as the noise level increases. This shows that our algorithm performs much better than the other algorithms in the presence of corruptions in the similarities.

The third row shows the results when Gaussian noise is added to the attributes. We observe that when the noise level increases, the errors of all algorithms increase. The errors of our algorithm do not increase much: they remain the best or close to the best up to the noise level 0.2 on all the data sets. Ward’s method also shows robustness, since the minimum variance criterion used is insensitive to this type of noise. The other algorithms generally show higher errors than our algorithms and Ward’s method.

In conclusion, our algorithm generally outperforms the other algorithms when corruption noise is added to the data attributes or the similarities, or when Gaussian noise is added to the data attributes. Its robustness to Gaussian noise in similarities is not so significant since such noise with large variance can change the neighbor rankings of all points considerably. Still, it can tolerate such noise when the noise variance is not too large.

Robustness to Parameter Tuning Our algorithm requires extra input parameters α and ν . There may be indirect ways to set their values, for example, by estimating the size of the smallest target cluster. Still, we are not aware of any efficient algorithm to compute the exact values. Here we show the robustness of our algorithm to parameter tuning. Note that the two parameters are always used together as the additive term $(\alpha + \nu)$, thus essentially the algorithm takes one parameter. So for evaluation, we vary the parameter $(\alpha + \nu)$




Figure 14: Classification Error of different algorithms in the inductive setting. The y -axis represents the % error. The x -axis represents data sets, where the numbers before the names of the data sets denote the fraction of similarities used by the inductive algorithms.

linearly and run our algorithm over these values.

Figure 13 shows the performance of the algorithm for different parameter values. We observe that the algorithm does not require the exact value of $(\alpha + \nu)$ as it shows good performance over a continuous range of values. The range is sufficiently large for all data sets except Iris. The range for Iris is relatively small as there is little noise in it, and thus the parameter should be set to small values. Generally, it is easy to land in the right range with only a few runs.

6.2.2 Inductive Setting

In this subsection, we present the evaluation results in the inductive setting. In this setting, the algorithm generates a hierarchy on a small random sample of the data set, and inserts the remaining points to generate a new hierarchy over the entire data set. We repeat the sampling and evaluation for 5 times and report the average results.

We compare our inductive algorithm with the random sample algorithm in (Eriksson, 2012). These algorithms sample some fraction of the similarities and use only these similarities. The percentage of sampled similarities can be tuned in these algorithms, so we compare their performance when they use the same amount of sampled similarities.

Figure 14 shows the results for eight configurations (using 5% or 10% similarities on four different data sets). Our algorithm consistently outperforms the random sampling algorithm.

Figure 15 shows the results on PFAM1 to PFAM10, which approximately satisfy the good neighborhood property (Voevodski et al., 2012). On all PFAM data sets, the errors of our algorithm are low while those of the random sample algorithm are much higher. This shows the significant advantage of our algorithm when the data approximately satisfies the good neighborhood property.




Figure 15: Classification Error on PFAM1 to PFAM10 data sets using 2.5% similarities. The y -axis in each case represents the % error, and the x -axis represents data sets.

7 Discussion

In this work we propose and analyze a new robust algorithm for bottom-up agglomerative clustering. We show that our algorithm can be used to cluster accurately in cases where the data satisfies a number of natural properties and where the traditional agglomerative algorithms fail. In particular, if the data satisfies the good neighborhood properties, the algorithm will be successful in generating a hierarchy such that the target clustering is close to a pruning of that hierarchy.

We also show how to extend our algorithm to the inductive setting, where the given data is only a small random sample of the entire data set. Our algorithm achieves similar correctness guarantees, requiring only a small random sample whose size is independent of that of the entire data set.

We empirically show that our algorithm consistently performs better than other hierarchical algorithms and are more robust to noise in the data. We also show the efficacy of the inductive version of our algorithm as a quicker alternative when evaluation over the complete data is resource intensive.

Additionally, our subsequent work (Balcan and Liang, 2013) showed that the algorithm can be applied to the closely related community detection task and compares favorably with existing approaches.

Acknowledgments: We thank Avrim Blum for numerous useful discussions, and Konstantin Voevodski for providing us the PFAM data sets.

References

- S. F. Altschul, W. Gish, W. Miller, E. W. Myers, and D. J. Lipman. Basic local alignment search tool. *Journal of Molecular Biology*, 1990.
- K. Bache and M. Lichman. UCI machine learning repository, 2013. URL <http://archive.ics.uci.edu/ml>.
- M. F. Balcan and Y. Liang. Modeling and detecting community hierarchies. In *Proceedings of the International Workshop on Similarity-Based Pattern Analysis and Recognition*, 2013.

- M. F. Balcan, A. Blum, and S. Vempala. A discriminative framework for clustering via similarity functions. In *Proceedings of the Annual ACM symposium on Theory of Computing*, 2008.
- M. F. Balcan, A. Blum, and A. Gupta. Clustering under approximation stability. *Journal of ACM*, 2013.
- A. Blum, A. Coja-Oghlan, A. Frieze, and S. Zhou. Separating populations with wide data: A spectral analysis. In *Algorithms and Computation*. 2007.
- D. Bryant and V. Berry. A structured family of clustering and tree construction methods. *Advances in Applied Mathematics*, 2001.
- K. Chaudhuri and S. Dasgupta. Rates of convergence for the cluster tree. *Advances in Neural Information Processing Systems*, 2010.
- D. Cheng, R. Kannan, S. Vempala, and G. Wang. A divide-and-merge methodology for clustering. *ACM Transaction on Database Systems*, 2006.
- S. Dasgupta and P. Long. Performance guarantees for hierarchical clustering. *Journal of Computer and System Sciences*, 2005.
- R.O. Duda, P.E. Hart, and D.G. Stork. *Pattern classification*. 2000.
- B. Eriksson. Hierarchical clustering using randomly selected measurements. In *Proceedings of the IEEE Statistical Signal Processing Workshop*, 2012.
- B. S. Everitt, S. Landau, M. Leese, and D. Stahl. *Hierarchical Clustering*. 2011.
- J. Feldman, R. O'Donnell, and R. A. Servedio. Learning mixtures of product distributions over discrete domains. *SIAM Journal on Computing*, 2008.
- S. Gollapudi, R. Kumar, and D. Sivakumar. Programmable clustering. In *Symposium on Principles of Database Systems*, 2006.
- J. C. Gower. A comparison of some methods of cluster analysis. *Biometrics*, 1967.
- S. Guha, R. Rastogi, and K. Shim. CURE: an efficient clustering algorithm for large databases. In *ACM SIGMOD Record*, 1998.
- A. K. Jain and R. C. Dubes. Algorithms for clustering data, 1981.
- A.K. Jain, M.N. Murty, and P.J. Flynn. Data clustering: a review. *ACM computing surveys*, 1999.
- S. C. Johnson. Hierarchical clustering schemes. *Psychometrika*, 1967.
- B. King. Step-wise clustering procedures. *Journal of the American Statistical Association*, 1967.
- H. W. Kuhn. The Hungarian method for the assignment algorithm. *Naval Research Logistics Quarterly*, 1955.
- Y. LeCun, L. Bottou, Y. Bengio, and P. Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 1998.

- M. Meilă and D. Heckerman. An experimental comparison of model-based clustering methods. *Machine Learning*, 2001.
- A. Moitra and M. Saks. A polynomial time algorithm for lossy population recovery. *arXiv preprint arXiv:1302.1515*, 2013.
- G. Nagy. State of the art in pattern recognition. *Proceedings of the IEEE*, 1968.
- M. Narasimhan, N. Jojic, and J. Bilmes. Q-clustering. *Advances in Neural Information Processing Systems*, 2006.
- M. Punta, P. C. Coghill, R. Y. Eberhardt, J. Mistry, J. Tate, C. Boursnell, N. Pang, K. Forslund, G. Ceric, J. Clements, A. Heger, L. Holm, E. L. L. Sonnhammer, S. R. Eddy, A. Bateman, and R. D. Finn. The pfam protein families database. *Nucleic Acids Research*, 2012.
- P. H. A. Sneath and R. R. Sokal. *Numerical taxonomy. The principles and practice of numerical classification*. 1973.
- K. Voevodski, M. F. Balcan, H. Röglin, S.-H. Teng, and Y. Xia. Active clustering of biological sequences. *Journal of Machine Learning Research*, 2012.
- J. H. Ward. Hierarchical grouping to optimize an objective function. *Journal of the American Statistical Association*, 1963.
- A. Wigderson and A. Yehudayoff. Population recovery and partial identification. In *Proceedings of the IEEE Annual Symposium on Foundations of Computer Science*, 2012.
- D. Wishart. Mode analysis: a generalization of nearest neighbour which reduces chaining effects. *Numerical Taxonomy*, 1969.

A Implementation Details of Algorithm 1

Here we give the full details of the implementation of Algorithm 1.

First, we need some auxiliary data structures to build the graphs F_t and H_t . See Algorithm 3 for the definitions of these data structures.

Second, we specify the order of merging clusters. In the merge step in Algorithm 1, the blobs in a sufficiently large connected component of H_t can be merged in arbitrary order. In our implementation, we merge two connected C_u, C_v in H_t such that they are not both singleton blobs and they have maximum $\text{median}_{x \in C_u, y \in C_v} S_t(x, y) / (|C_u| + |C_v|)$ (so that we are most confident about merging them). Then we merge singleton clusters.

Third, for practical purposes, we can slightly modify the algorithm to speed it up on practical instances⁵. When there are less than $4(\alpha + \nu)n$ singleton blobs, we know that they cannot be merged together into one blob. So we can simply merge each singleton blob with the non-singleton blob that has the highest median similarity. This will correctly assign all but bad points under the good neighborhood properties. Similarly, when the number of singleton blobs is less than half the current threshold, we can safely merge each singleton blob with the non-singleton blob that has the highest median similarity.

⁵This does not change the time complexity and the correctness, but we observe that it helps speed up practical instances.

Algorithm 3 Implementation Details of Robust Median Neighborhood Linkage

Input: Similarity function \mathcal{K} on a set of points S , $n = |S|$, $\alpha > 0, \nu > 0$.

- Step 1** For each point, sort the other points increasingly according to the distances.
Initialize $t = 6(\alpha + \nu)n + 1$, and \mathcal{C}'_{t-1} to be a list of singleton blobs.
while $|\mathcal{C}'_{t-1}| > 1$ **do**
- Step 2** \triangleright Build a graph F_t on the points in S as follows:
Set $I_t(x, y) = 1$ if y is in x 's t nearest neighbors; $I_t(x, y) = 0$ otherwise.
Set $N_t = I_t(I_t)^T$.
Set $F_t(x, y) = 1$ if $N_t(x, y) \geq t - 2(\alpha + \nu)n$; $F_t(x, y) = 0$ otherwise.
- Step 3** \triangleright Build a graph H_t on the blobs in \mathcal{C}'_{t-1} as follows:
Set $NS_t = F_t(F_t)^T$.
Set $FC_t(x, y) = 1$ if x, y are in the same blob in \mathcal{C}'_{t-1} and $F_t(x, y) = 1$;
 $FC_t(x, y) = 0$ otherwise.
Set $S_t = F_t(FC_t)^T + FC_t(F_t)^T$.
for any $C_u, C_v \in \mathcal{C}'_{t-1}$ **do**
 if $C_u = \{x\}$ and $C_v = \{y\}$ **then**
 $H_t(C_u, C_v) = 1$ if $NS_t(x, y) > (\alpha + \nu)n$;
 $H_t(C_u, C_v) = 0$ otherwise.
 else
 $H_t(C_u, C_v) = 1$ if $\text{median}_{x \in C_u, y \in C_v} S_t(x, y) > \frac{|C_u| + |C_v|}{4}$;
 $H_t(C_u, C_v) = 0$ otherwise.
 end if
end for
- Step 4** \triangleright Merge blobs
while $\exists C_u, C_v$ with $H_t(C_u, C_v) = 1$ and $|C_u| + |C_v| > 4(\alpha + \nu)n$ **do**
 Find the pair C_u, C_v with maximum $\text{median}_{x \in C_u, y \in C_v} \frac{S_t(x, y)}{|C_u| + |C_v|}$.
 Merge the pair C_u, C_v , and update \mathcal{C}'_{t-1} . Recompute FC_t, S_t and H_t .
end while
- Step 5** \triangleright Merge singletons
while \exists component V in H_t with $|\cup_{C \in V} C| \geq 4(\alpha + \nu)n$ **do**
 Merge blobs in V , and update \mathcal{C}'_{t-1} . Recompute FC_t, S_t and H_t .
end while
- Step 6** \triangleright Speed up
if \exists less than $\max\{4(\alpha + \nu)n, t/2\}$ singleton blobs in \mathcal{C}'_{t-1} **then**
 Merge each singleton with the non-singleton blob of highest median similarity.
 Update \mathcal{C}'_{t-1} . Recompute FC_t, S_t and H_t .
end if
 $\mathcal{C}'_t = \mathcal{C}'_{t-1}$.
- Step 7** \triangleright Increase threshold
 $t = t + 1$.
end while

Output: Tree T with single points as leaves and internal nodes corresponding to the merges performed.

B Additional Proofs for Section 5

Here we provide the details for proving that Algorithm 2 also succeeds for weak good neighborhood. First, by a similar argument as that in Lemma 6, we can prove Lemma 9 showing that for a fixed p in $X \setminus B$ and fixed $x \in A_p$, the first condition of the weak good neighborhood is still satisfied on a sufficiently large sample (Recall the definition of Property 5). Similarly, we can prove Lemma 10 showing that the second condition of the weak good neighborhood is also satisfied. Then, the similarity \mathcal{K} satisfies weak good neighborhood property with respect to the clustering induced over the sample (Lemma 8). Our final guarantee, Theorem 4, then follows from the lemmas.

Lemma 9. *Let \mathcal{K} be a symmetric similarity function satisfying the weak (α, β, ν) -good neighborhood for the clustering problem (X, ℓ) . Consider any fixed $p \in X \setminus B$ and any fixed $x \in A_p$. If the size of the sample satisfies $n = \Theta\left(\frac{1}{\alpha} \ln \frac{1}{\delta}\right)$, then with probability at least $1 - \delta$, x has at most $2\alpha n$ neighbors outside $A_p \cap S$ out of the $|A_p \cap S|$ nearest neighbors in $S \setminus B$.*

Lemma 10. *Let \mathcal{K} be a symmetric similarity function satisfying the weak (α, β, ν) -good neighborhood for the clustering problem (X, ℓ) . Consider any fixed $p \in X \setminus B$ and any fixed good point $x \in A_p$. If the size of the sample satisfies $n = \Theta\left(\frac{1}{\alpha} \ln \frac{1}{\delta}\right)$, then with probability at least $1 - \delta$, x has at most $2\alpha n$ neighbors outside $C(x) \cap S$ out of the $|A_p \cap S|$ nearest neighbors in $S \setminus B$.*

Lemma 8 *Let \mathcal{K} be a symmetric similarity function satisfying the weak (α, β, ν) -good neighborhood for the clustering problem (X, ℓ) . Furthermore, it satisfies that for any $p \in X \setminus B$, $|A_p| > 24(\alpha + \nu)N$. If the size of the sample satisfies $n = \Theta\left(\frac{1}{\min(\alpha, \nu)} \ln \frac{1}{\delta \min(\alpha, \nu)}\right)$, then with probability at least $1 - \delta$, \mathcal{K} satisfies the $(2\alpha, \frac{15}{16}\beta, 2\nu)$ -good neighborhood with respect to the clustering induced over the sample S .*

Proof. Consider the first condition of the weak good neighborhood property. First, by Chernoff bounds, when $n \geq \frac{3}{\nu} \ln \frac{4}{\delta}$, we have that with probability at least $1 - \delta/4$, at most $2\nu n$ bad points fall into the sample. Next, by Lemma 9 and union bound, when $n = \Theta\left(\frac{1}{\alpha} \ln \frac{1}{\delta}\right)$ we have that with probability at least $1 - \delta/4$, for any point $p \in S \setminus B$, any point $x \in A_p \cap S$ has at most $2\alpha n$ neighbors outside $A_p \cap S$ out of the $|A_p \cap S|$ nearest neighbors in $S \setminus B$. Since $|A_p| > 24(\alpha + \nu)N$, we also have $|A_p \cap S| > 12(\alpha + \nu)n$ with probability at least $1 - \delta/4$. So the first condition of the weak good neighborhood property is satisfied.

Now consider the second condition. Fix C_i and a point $p \in (C_i \setminus B) \cap S$. When $n = \Theta\left(\frac{1}{\alpha} \ln \frac{n}{\delta}\right)$, with probability at least $1 - \delta/(8n)$, at least $\frac{15}{16}\beta$ fraction of points x in $A_p \cap S$ have all but at most αN nearest neighbors from $C_i \setminus B$ out of their $|C_i \setminus B|$ nearest neighbors in $X \setminus B$. Fix such a point $x \in A_p \cap S$. By Lemma 10, with probability at least $1 - \delta/(8n^2)$, it has all but at most $2\alpha n$ nearest neighbors from $(C_i \setminus B) \cap S$ out of their $|(C_i \setminus B) \cap S|$ nearest neighbors in $S \setminus B$. By union bound, we have that with probability at least $1 - \delta/4$, for any C_i and any $p \in C_i \setminus B$, at least $\frac{15}{16}\beta$ fraction of points in $A_p \cap S$ have all but at most $2\alpha n$ nearest neighbors from $(C_i \setminus B) \cap S$ out of their $|(C_i \setminus B) \cap S|$ nearest neighbors in $S \setminus B$. So the second condition is also satisfied.

Therefore, if $n = \Theta\left(\frac{1}{\min(\alpha, \nu)} \ln \frac{n}{\delta}\right)$, then with probability at least $1 - \delta$, the similarity function satisfies the $(2\alpha, 2\nu)$ -good neighborhood property with respect to the clustering induced over the sample S . The lemma then follows from the fact that $n = \Theta\left(\frac{1}{\min(\alpha, \nu)} \ln \frac{1}{\delta \min(\alpha, \nu)}\right)$ implies $n = \Theta\left(\frac{1}{\min(\alpha, \nu)} \ln \frac{n}{\delta}\right)$. \square

Theorem 4 *Let \mathcal{K} be a symmetric similarity function satisfying the weak (α, β, ν) -good neighborhood for the clustering problem (X, ℓ) with $\beta \geq \frac{14}{15}$. Furthermore, it satisfies that for any $p \in X \setminus B$, $|A_p| > 24(\alpha + \nu)N$. Then Algorithm 2 with parameters $n = \Theta\left(\frac{1}{\min(\alpha, \nu)} \ln \frac{1}{\delta \min(\alpha, \nu)}\right)$ produces a hierarchy with a pruning that is $(\nu + \delta)$ -close to the target clustering with probability $1 - \delta$.*

Proof. Note that by Lemma 8, with probability at least $1 - \delta/4$, we have that \mathcal{K} satisfies the weak $(2\alpha, \frac{15}{16}\beta, 2\nu)$ -good neighborhood with respect to the clustering induced over the sample. Then by Theorem 1, Algorithm 1 outputs a hierarchy T on the sample S with a pruning $\{u_1, \dots, u_k\}$ such that $u_i \setminus B = (C_i \cap S) \setminus B$.

Now we want to show that f_{u_1}, \dots, f_{u_k} have error at most $\nu + \delta$ with probability at least $1 - \delta/2$. For convenience, let $u(x)$ be a shorthand of $u_{\ell(x)}$. Then it is sufficient to show that with probability at least $1 - \delta/2$, a $(1 - \delta)$ fraction of points $x \in X \setminus B$ have $f_{u(x)}(x) = 1$. Fix C_i and a point $x \in C_i \setminus B$. By Lemma 9, with probability at least $1 - \delta^2/2$, out of the $|A_x \cap S|$ nearest neighbors of x in $S \setminus B$, at most $2\alpha n$ can be outside A_x . Then out of the $6(\alpha + \nu)n$ nearest neighbors of x in S , at most $2(\alpha + \nu)n$ points are outside $A_x \cap S$. By Lemma 1, u_i contains $A_x \cap S$, so u_i must contain at least $4(\alpha + \nu)n$ points in $N_S(x)$. Consequently, any ancestor w of u_i , including u_i , has more points in $N_S(x)$ than any other sibling of w . Then we must have $f_w(x) = 1$ for any ancestor w of u_i . In particular, $f_{u_i}(x) = 1$. So, for any point $x \in X \setminus B$, with probability at least $1 - \delta^2/2$ over the draw of the random sample, $f_{u(x)}(x) = 1$. By Markov inequality, with probability at least $1 - \delta/2$, a $(1 - \delta)$ fraction of points $x \in X \setminus B$ have $f_{u(x)}(x) = 1$. \square